

# Introduction to DelftBlue

Dennis Palagin  
20 September 2024

---

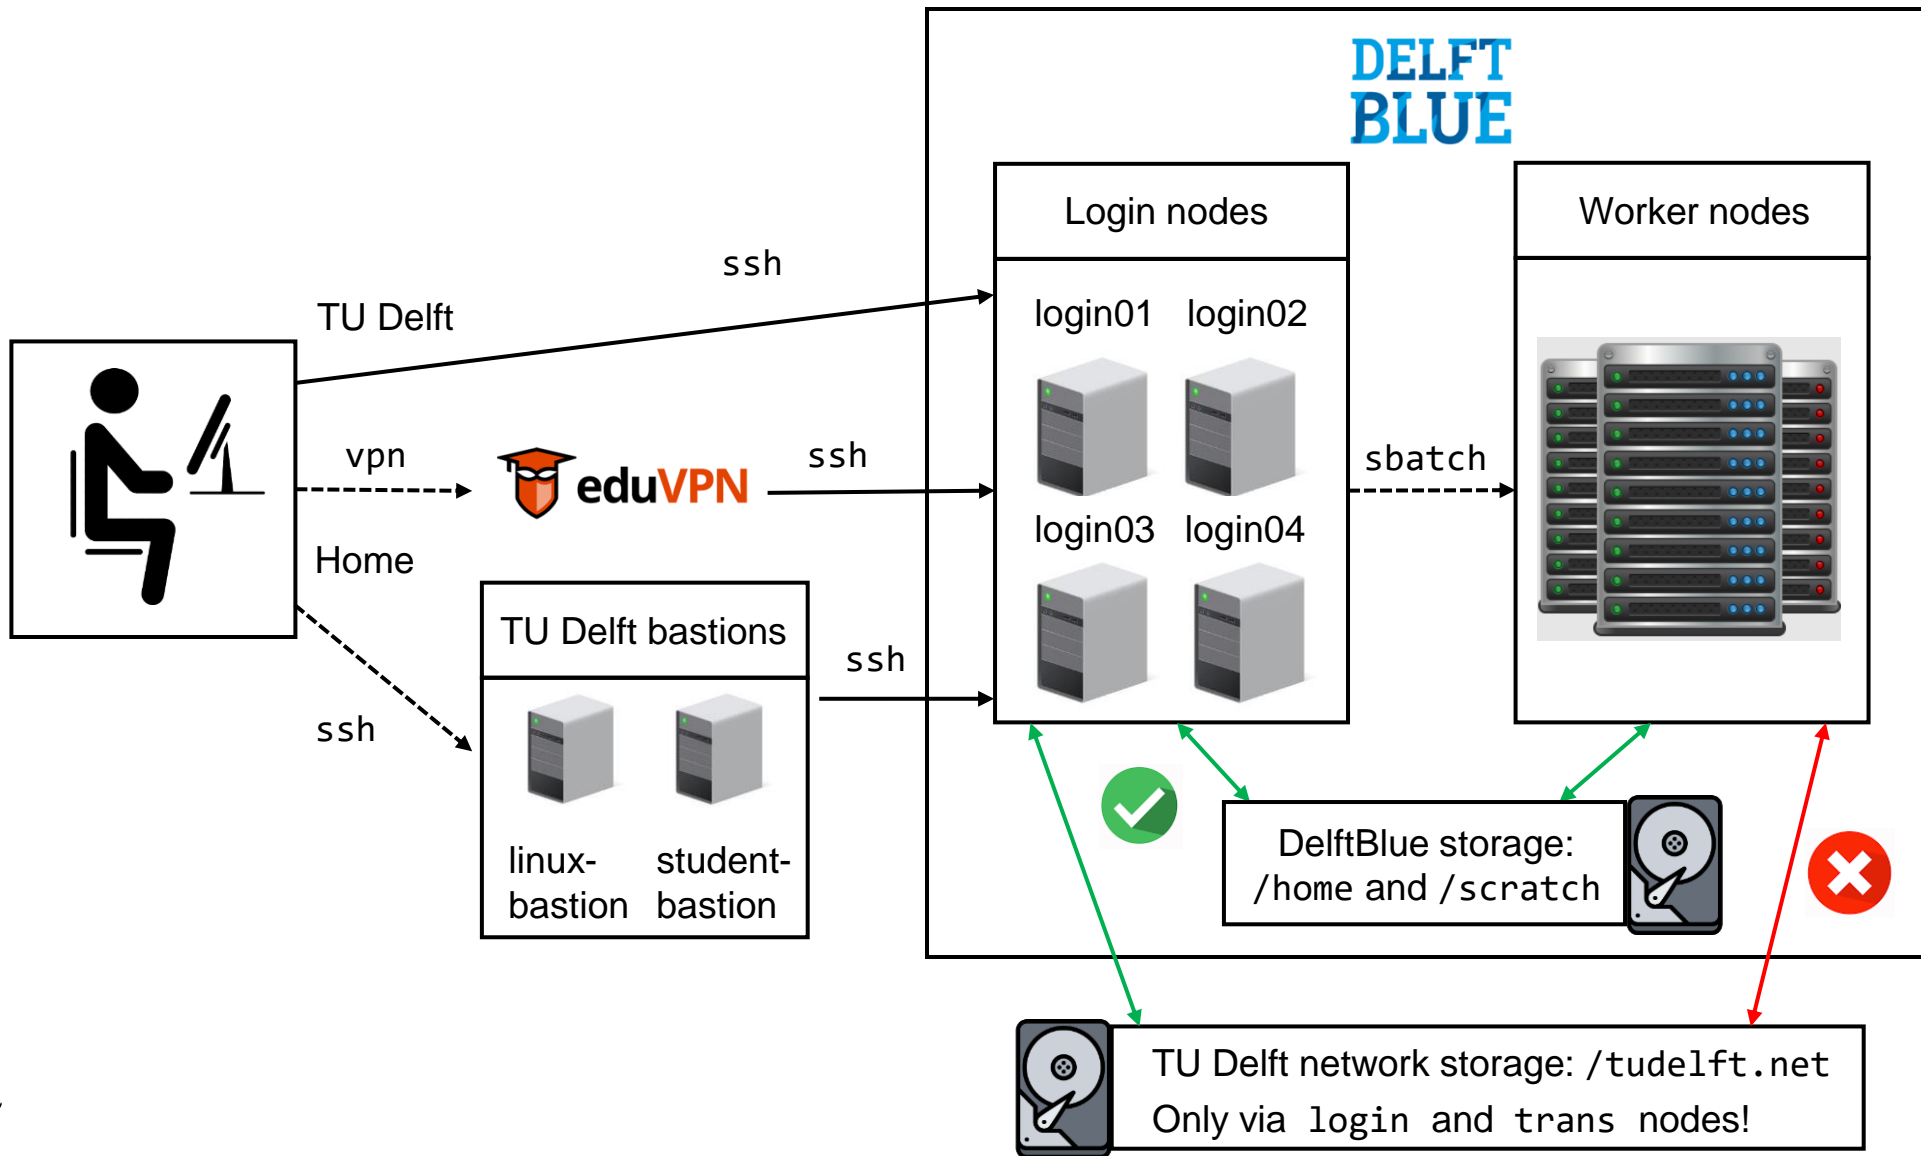


# Practical use of DelftBlue

- Recap: what is a cluster computer
- DHPC and DelftBlue: organization
- Accessing the system
- File systems and data transfer
- Queuing, accounting
- Module system (Imod)

# Practical use of DelftBlue

- Recap: what is a cluster computer
- DHPC and DelftBlue: organization
- Accessing the system
- File systems and data transfer
- Queuing, accounting
- Module system (Imod)





# DelftBlue

- Fast and flexible
- 16.000 CPU cores
- 20 GPU nodes
- High-Speed Interconnect based on Mellanox InfiniBand
- And a 700TB high-speed parallel storage subsystem

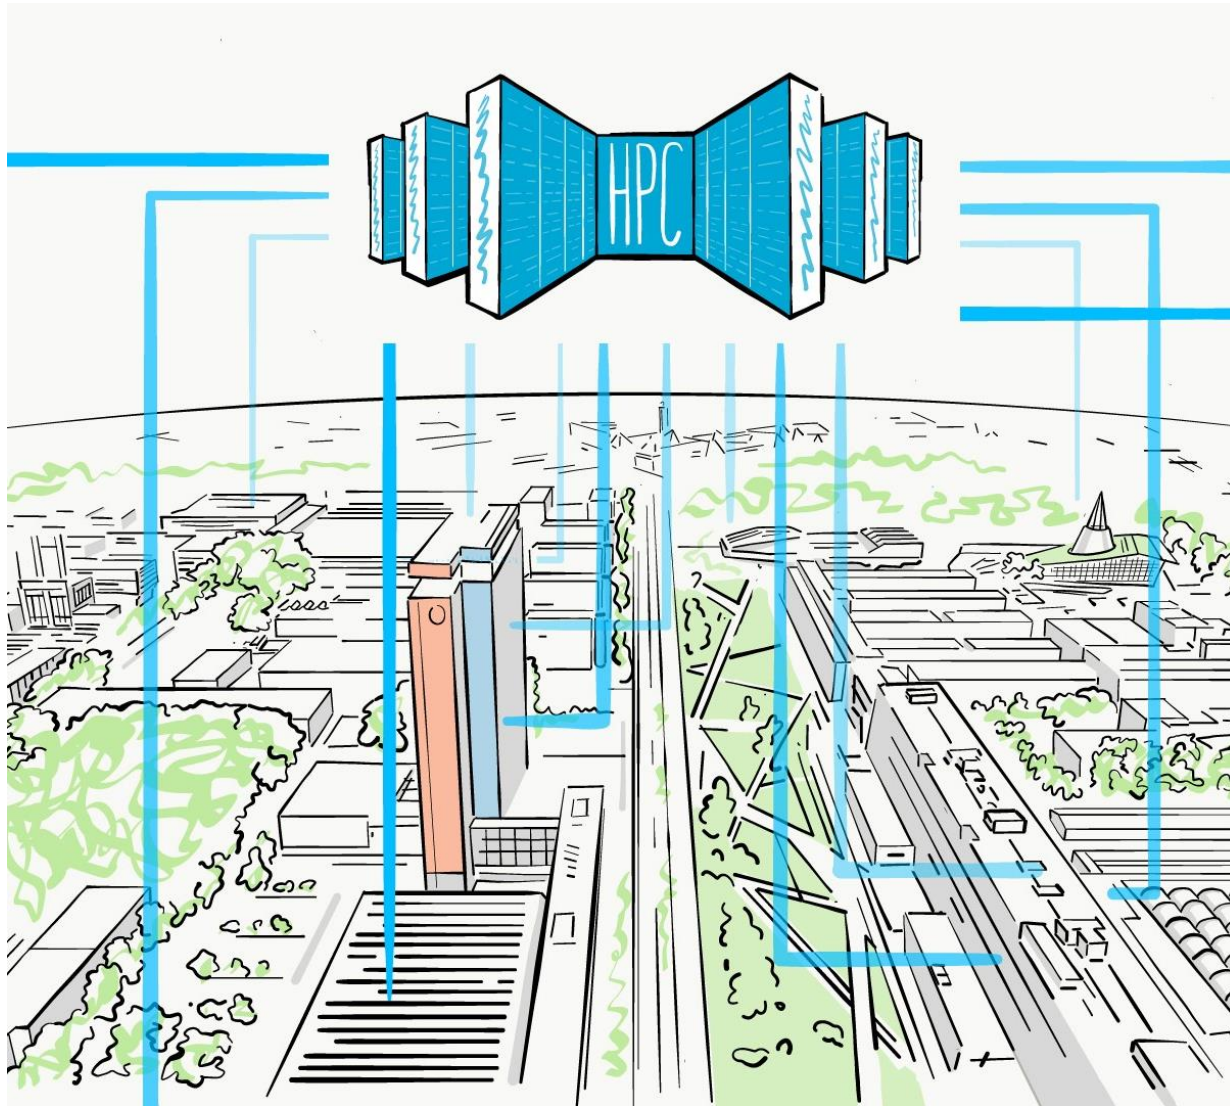


# Practical use of DelftBlue

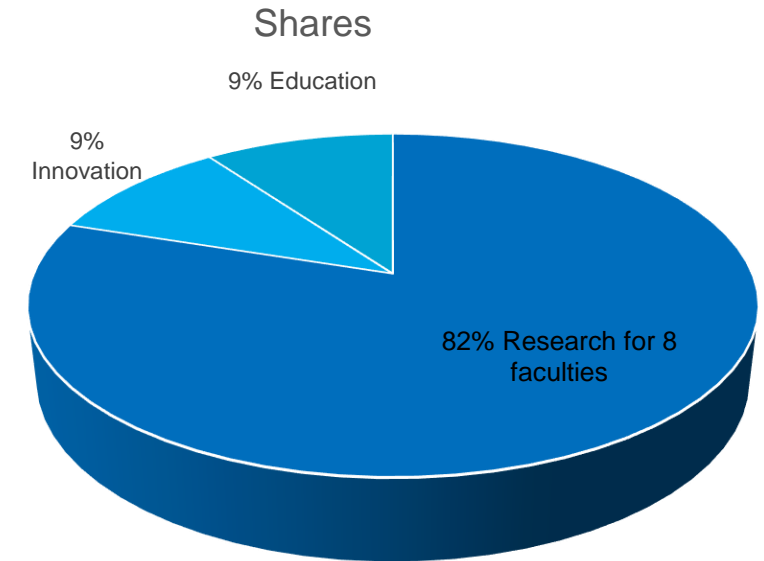
- Recap: what is a cluster computer
- **DHPC and DelftBlue: organization**
- Accessing the system
- File systems and data transfer
- Queuing, accounting
- Module system (Imod)



# DelftBlue community: 3mE, ABE, AE, AS, CEG, EEMCS, IDE, TPM, QuTech



**TU**Delft



Research AE/LR	12%
Research EEMCS/EWI	12%
Research TPM/TBM	12%
Research CEG/CiTG	12%
Research 3mE	12%
Research AS/TNW	12%
Research ABE/BK	6%
Research IDE/IO	2%
Research Qutech	2%
Education	9%
Innovation	9%
	100%

# Documentation:

- Documentation:

<https://www.tudelft.nl/dhpc/documentation>

- Mattermost:

<https://mattermost.tudelft.nl/dhpc/>

- Self Service Portal (TopDesk):

<https://tudelft.topdesk.net>

- Training and courses:

<https://www.tudelft.nl/cse/education/courses>



# Practical use of DelftBlue

- Recap: what is a cluster computer
- DHPC and DelftBlue: organization
- **Accessing the system**
- File systems and data transfer
- Queuing, accounting
- Module system (Imod)

# ssh

Anyone with TU Delft <netid> should be able to SSH to DelftBlue:

```
user@laptop:~ $ ssh <netid>@login.delftblue.tudelft.nl
```

This will log you in into one of the four login nodes (login01, login02, login03, or login04).  
Your home directory is directly accessible (via /home/<netid>/):

```
[<netid>@login01 ~]$ echo $HOME  
/home/NetID
```

```
[<netid>@login01 ~]$ ls -l
```

```
drwxr-xr-x  3 <netid> domain users      1 Mar 17 13:51 calcs  
drwxr-xr-x  3 <netid> domain users      1 Mar 17 13:27 codes  
lrwxrwxrwx  1 <netid> domain users    21 Mar 17 14:27 scratch -> /scratch/NetID  
drwxr-xr-x  3 <netid> domain users      1 Mar 17 13:27 tools
```

# Graphical tools Windows: PuTTY, Bitvise SSH, MobaXterm, etc...

The screenshot displays a Windows desktop environment with several graphical tools. The primary focus is on the Bitvise SSH client, which is running a terminal session. The terminal output shows a file listing command being executed, displaying a directory structure with permissions, owners, sizes, and dates. The desktop also features icons for Remote Desktop, Notepad++, kdiff3, and Git Bash. A Bitvise SFTP window is also open, showing a file browser interface with local and remote file lists.

**Bitvise SSH Client - Terminal Output:**

```
delftblue.nl.tlp - dpalagin@login.delftblue.tudelft.nl:22 - Bitvise xterm - dpalagin@login02:~  
[dpalagin@login02 ~]$ ls -la  
total 41  
drwx----- 9 dpalagin domain users 16 Mar 18 11:18 .  
drwxr-xr-x 16 root root 14 Mar 17 22:34 ..  
-rw----- 1 dpalagin domain users 7777 Mar 18 11:35 .bash_history  
-rw----- 1 dpalagin domain users 18 Mar 17 10:57 .bash_logout  
-rw----- 1 dpalagin domain users 141 Mar 17 10:57 .bash_profile  
-rw----- 1 dpalagin domain users 656 Mar 17 14:08 .bashrc  
drwxr-xr-x 3 dpalagin domain users 1 Mar 17 13:51 calcs  
drwxr-xr-x 3 dpalagin domain users 1 Mar 17 13:27 codes  
drwx----- 3 dpalagin domain users 1 Mar 17 13:28 .config  
-rw----- 1 dpalagin domain users 334 Mar 17 10:57 .emacs  
-rw----- 1 dpalagin domain users 172 Mar 17 10:57 .kshrc  
drwxr-xr-x 3 dpalagin domain users 1 Mar 17 14:29 .ondemand  
lrwxrwxrwx 1 dpalagin domain users 21 Mar 17 14:27 scratch -> /mnt/scratch/dpalagin  
drwx----- 2 dpalagin domain users 4 Mar 17 11:10 .ssh  
drwxr-xr-x 3 dpalagin domain users 1 Mar 17 13:27 tools  
drwxr-xr-x 2 dpalagin domain users 1 Mar 17 14:28 .vim  
-rw----- 1 dpalagin domain users 24969 Mar 18 11:18 .viminfo  
-rw-r--r-- 1 dpalagin domain users 165 Mar 17 13:27 .wget-hsts  
[dpalagin@login02 ~]$
```

**Bitvise SFTP Client - Remote Files:**

Name	Size	Type	Date Modified	Permissions
.config	1	File folder	17/03/2022 13:28	drw
.ssh	4	File folder	17/03/2022 11:10	drw
.vim	1	File folder	17/03/2022 14:28	drw
calcs	1	File folder	17/03/2022 13:51	drw
codes	1	File folder	17/03/2022 13:27	drw
ondemand	1	File folder	17/03/2022 14:29	drw
scratch	21	File folder	17/03/2022 14:27	lws
tools	1	File folder	17/03/2022 13:27	drw
.bash_history	7,777	BASH_HIS...	18/03/2022 11:35	-rw
.bash_logout	18	BASH_LOG...	17/03/2022 10:57	-rw
.bash_profile	141	BASH_PRO...	17/03/2022 10:57	-rw
.bashrc	656	BASHRC File	17/03/2022 14:08	-rw
.emacs	334	EMACS File	17/03/2022 10:57	-rw
.kshrc	172	KSHRC File	17/03/2022 10:57	-rw
.viminfo	24,969	VIMINFO File	18/03/2022 11:18	-rw
.wget-hsts	165	WGGET-HST...	17/03/2022 13:27	-rw



# Web tools: OpenOnDemand

← → ↻ https://login.delftblue.tudelft.nl/pun/sys/dashboard/files/fs/home/dpalagin ☆ 🔔 ☰

Open OnDemand Files ▾ Jobs ▾ Clusters ▾ Interactive Apps ▾ 📁

🔍 ? 👤 ➦

📄 Open in Terminal ➤ ➕ New File 📁 New Directory 📶 Upload 📶 Download 📄 Copy/Move 🗑 Delete

Home Directory

📶 / home / dpalagin / 📄 Change directory 📄 Copy path

☐ Show Owner/Mode ☐ Show Dotfiles Filter:

Showing 5 of 16 rows - 0 rows selected

	Type	↑ ↓ Name	↑ ↓	Size	↑ ↓ Modified at	↑ ↓
<input type="checkbox"/>	📁	calcs	⋮ ▾	-	17/03/2022 13:51:50	
<input type="checkbox"/>	📁	codes	⋮ ▾	-	17/03/2022 13:27:10	
<input type="checkbox"/>	📁	ondemand	⋮ ▾	-	17/03/2022 14:29:46	
<input type="checkbox"/>	📁	scratch	⋮ ▾	-	17/03/2022 14:28:05	
<input type="checkbox"/>	📁	tools	⋮ ▾	-	17/03/2022 13:27:19	

powered by **OPEN** **OnDemand**

OnDemand version: v2.0.20

# Practical use of DelftBlue

- Recap: what is a cluster computer
- DHPC and DelftBlue: organization
- Accessing the system
- **File systems and data transfer**
- Queuing, accounting
- Module system (Imod)

# scp

1. Transfer a file from your computer to DelftBlue:

```
user@laptop:~ $ scp localfile <netid>@login.delftblue.tudelft.nl:~/destination_on_DelftBlue/
```

2. Transfer a folder from your computer to DelftBlue:

```
user@laptop:~ $ scp -r localfile <netid>@login.delftblue.tudelft.nl:~/destination_on_DelftBlue/
```

3. Transfer a file from DelftBlue to your computer:

```
user@laptop:~ $ scp <netid>@login.delftblue.tudelft.nl:~/folder_on_DelftBlue/remotefile ./
```

4. Transfer a folder from DelftBlue to your computer:

```
user@laptop:~ $ scp -r <netid>@login.delftblue.tudelft.nl:~/folder_on_DelftBlue ./
```



# Network drives

```
[NetID@login02 tudelft.net]$ ls -l
```

```
total 104
```

drwxr-xr-x	12	root	root	4096	Mar	16	15:16	staff-bulk
drwxr-xr-x	13	root	root	4096	Mar	16	15:16	staff-groups
drwxr-xr-x	28	root	root	4096	Mar	17	11:18	staff-homes
drwxr-xr-x	28	root	root	4096	Mar	17	11:18	staff-homes-linux
drwxr-xr-x	1674	root	root	65536	Mar	17	11:18	staff-umbrella
drwxr-xr-x	12	root	root	4096	Mar	17	11:18	student-groups
drwxr-xr-x	28	root	root	4096	Mar	17	11:18	student-homes
drwxr-xr-x	28	root	root	4096	Mar	17	11:18	student-homes-linux

# Exercise 101

- ssh to DelftBlue

```
user@laptop:~ $ ssh NetID@login.delftblue.tudelft.nl
```

- Download exercises from <https://doc.dhpc.tudelft.nl/delftblue/>
- Copy to DelftBlue and unzip.

# Practical use of DelftBlue

- Recap: what is a cluster computer
- DHPC and DelftBlue: organization
- Accessing the system
- File systems and data transfer
- **Queuing, accounting**
- Module system (Imod)



# What is a scheduler?

```

34101180 kchoudhu pr89 5vb2_ctd_unfol PD Priority Tomorr 20:15 0:00 1-00:00:00 16 192
34101393 kchoudhu pr89 5vb2_ctd_unfol PD Priority Tomorr 20:15 0:00 1-00:00:00 16 192
34101545 ykarami pr118 IgG1C1 PD Priority Tomorr 20:15 0:00 1-00:00:00 6 72
34101588 ykarami pr118 IgG3-1C1 PD Priority Tomorr 20:15 0:00 1-00:00:00 6 72
34101854 mperisdi pr107 Debug PD Priority Tomorr 20:15 0:00 1-00:00:00 3 3
34102044 calleva pr89 4q4g PD Priority Tomorr 20:15 0:00 1-00:00:00 8 96
34102045 calleva pr89 4g4q PD Priority Tomorr 20:15 0:00 1-00:00:00 8 96
34102780 asridhar pr89 Kv3_WT PD Priority Tomorr 20:15 0:00 1-00:00:00 12 144
34102782 asridhar pr89 Kv3_WT_2 PD Priority Tomorr 20:15 0:00 1-00:00:00 12 144
34102783 asridhar pr89 D120_V253 PD Priority Tomorr 20:15 0:00 1-00:00:00 8 96
34102784 asridhar pr89 D120_V253_2 PD Priority Tomorr 20:15 0:00 1-00:00:00 8 96
34102785 asridhar pr89 F256A PD Priority Tomorr 20:15 0:00 1-00:00:00 8 96
34102786 asridhar pr89 F256A_2 PD Priority Tomorr 20:15 0:00 1-00:00:00 8 96
34102787 asridhar pr89 gaba5 PD Priority Tomorr 20:15 0:00 1-00:00:00 12 144
34102788 asridhar pr89 gaba4 PD Priority Tomorr 20:15 0:00 1-00:00:00 12 144
34102792 asridhar pr89 gaba6 PD Priority Tomorr 20:15 0:00 1-00:00:00 12 144
34102827 mbiliche pr66 Gkclp5_825 PD Priority Tomorr 20:15 0:00 1-00:00:00 25 300
34102868 mbiliche pr66 Gkclp4_866 PD Priority Tomorr 20:15 0:00 1-00:00:00 25 300
34102902 mbiliche pr66 Gkclp3_854 PD Priority Tomorr 20:15 0:00 1-00:00:00 25 300
34103693 kchoudhu pr89 5vb2_ctd_unfol PD Priority Tomorr 20:15 0:00 1-00:00:00 16 192
34105695 akumawat pr117 ICA-SOFT-AB2.1 PD Priority Tomorr 20:15 0:00 1-00:00:00 8 96
34105696 akumawat pr117 ICA-SOFT-AB2.2 PD Priority Tomorr 20:15 0:00 1-00:00:00 8 96
34105700 akumawat pr117 ICA-SOFT-AB2.3 PD Priority Tomorr 20:15 0:00 1-00:00:00 8 96
34105736 akumawat pr117 MUT-NSc-AB2.9 PD Priority Tomorr 20:15 0:00 1-00:00:00 8 96
34105739 akumawat pr117 MUT-NSc-AB2.10 PD Priority Tomorr 20:15 0:00 1-00:00:00 8 96
34105795 ykarami pr118 IgG3-2C1 PD Priority Tomorr 20:15 0:00 1-00:00:00 6 72
34105821 akumawat pr117 AMB19-AB2.1 PD Priority Tomorr 20:15 0:00 1-00:00:00 8 96
34108696 ykarami pr118 MTDP PD Priority Tomorr 20:15 0:00 1-00:00:00 100 1200
34108877 ykarami pr118 IgG3-1C2 PD Priority Tomorr 20:15 0:00 1-00:00:00 6 72
34111511 sfurini pr107 e4f PD Priority Tomorr 20:15 0:00 1-00:00:00 8 192
34087143 ykarami pr118 IgG1B3 CG None Ystday 12:28 20:54:35 3:05:25 1 144
dpalagin@daint105:~>

```

# How do I work with cluster?

1. **Prepare input files for your code** on a personal computer
2. **Upload input files** and required data to the cluster's storage
3. **Determine** required **resources**
4. **Create job script**
5. **Submit** job(s) to scheduler
6. **Monitor** progress (via output files) and resource use (via statistics)
7. **Download** results to personal computer for further processing
8. **Cleanup** files

# Typical commands?

- Create a job script in a file

```
#!/bin/sh

#SBATCH --job-name=job_name
#SBATCH --partition=compute
#SBATCH --account=research-eemcs-diam
#SBATCH --time=01:00:00
#SBATCH --ntasks=4
#SBATCH --cpus-per-task=1
#SBATCH --mem-per-cpu=1G

module load 2023r1
module load openmpi
srun ./executable > output.log
```

- system info: type of script
- select job name
- select partition to run your job on
- specify account
- request run time
- number of tasks (parallel)
- CPUs (threads) per task
- RAM per CPU
- Use DelftBlue software collection
- load openmpi module
- start tasks with srun



# Typical commands?

- Log in to one of the login nodes
- Submit job using the job script

```
$ sbatch jobscript.sbatch  
Submitted batch job 1
```

- See queue status

```
$ squeue  
JOBID PARTITION      NAME      USER ST  TIME  
  
1      general    jobscrip  somebody R 0:01
```

- See job output

```
$ cat slurm-1.out  
Hello world!
```

- Cancel job

```
$ scancel 1  
$ squeue  
JOBID PARTITION      NAME      USER ST  TIME  
NODES NODELIST(REASON)
```

# Example 1: Hello, World! on 4 CPUs

Our first submission script `helloworld.sh`:

```
#!/bin/bash
#SBATCH --job-name="01_hello"
#SBATCH --time=00:10:00
#SBATCH --ntasks=4
#SBATCH --cpus-per-task=1
#SBATCH --partition=compute
#SBATCH --mem=1GB
#SBATCH --account=Education-EEMCS-Courses-WI4049TU

echo "Hello, World!" >> helloworld.txt
echo "The following nodes are reporting for duty:" >> helloworld.txt
srun hostname >> helloworld.txt
echo "Have a great day!" >> helloworld.txt
```

# Example 1: Hello, World! on 4 CPUs

```
NetID@login01:~ $ sbatch helloworld.sh
```

```
Hello, World!  
The following nodes are reporting for duty:  
cmpXXX  
cmpXXX  
cmpXXX  
cmpXXX  
Have a great day!
```

## Example 1.2: Hello, World!

Our submission script `helloworld2nodes.sh`:

```
#!/bin/bash
#SBATCH --job-name="01_hello"
#SBATCH --time=00:10:00
#SBATCH --nodes=2
#SBATCH --ntasks-per-node=2
#SBATCH --cpus-per-task=1
#SBATCH --partition=compute
#SBATCH --mem=1GB
#SBATCH --account=Education-EEMCS-Courses-WI4049TU

echo "Hello, World!" >> helloworld.txt
echo "The following nodes are reporting for duty:" >> helloworld.txt
srun hostname >> helloworld.txt
echo "Have a great day!" >> helloworld.txt
```

## Example 1.2: Hello, World!

```
NetID@login01:~ $ sbatch helloworld2nodes.sh
```

```
Hello, World!  
The following nodes are reporting for duty:  
cmpXXX  
cmpXXX  
cmpYYY  
cmpYYY  
Have a great day!
```



# Practical use of DelftBlue

- Recap: what is a cluster computer
- DHPC and DelftBlue: organization
- Accessing the system
- File systems and data transfer
- Queuing, accounting
- Module system (Imod)

# Module system demo

```
[NetID@login02 ~]$ module avail
```

→ list available modules

```
[NetID@login02 ~]$ module load
```

→ load module

```
[NetID@login02 ~]$ module unload
```

→ unload module

```
[NetID@login02 ~]$ module list
```

→ list loaded modules

```
[NetID@login02 ~]$ module spider {module}
```

→ find module {module}

Trilinos example: <https://doc.dhpc.tudelft.nl/delftblue/DHPC-modules/>

## Example 2: Julia

Let's prepare a little program to draw a Mandelbrot set:

```
function mandelbrot(a)
    z = 0
    for i=1:50
        z = z^2 + a
    end
    return z
end

for y=1.0:-0.05:-1.0
    for x=-2.0:0.0315:0.5
        abs(mandelbrot(complex(x, y))) < 2 ? print("*")
    : print(" ")
    end
    println()
end
```

## Example 2: Julia

Submission script `run_julia_mandelbrot.sh`:

```
#!/bin/bash
#
#SBATCH --job-name="julia"
#SBATCH --time=00:10:00
#SBATCH --partition=compute
#SBATCH --ntasks=1
#SBATCH --cpus-per-task=1
#SBATCH --mem-per-cpu=1G
#SBATCH --account=Education-EEMCS-Courses-WI4049TU

module load 2023r1
module load julia

srun julia mandelbrot.jl
```

## Example 2: Julia

Let's submit the job and check out the result:

Let's submit the job and check out the result:

[illegible]

## Example 3: Python with MPI: mpi4py

Let's prepare a little program to calculate number pi in Python called `calculate_pi.py`:

```
from mpi4py import MPI
from math import pi as PI
from numpy import array

def comp_pi(n, myrank=0, nprocs=1):
    h = 1.0 / n
    s = 0.0
    for i in range(myrank + 1, n + 1, nprocs):
        x = h * (i - 0.5)
        s += 4.0 / (1.0 + x**2)
    return s * h

def prn_pi(pi, PI):
    message = "pi is approximately %.16f, error is  
%.16f"
    print (message % (pi, abs(pi - PI)))
```



## Example 3: Python with MPI: mpi4py

```
comm = MPI.COMM_WORLD
nprocs = comm.Get_size()
myrank = comm.Get_rank()

n = array(0, dtype=int)
pi = array(0, dtype=float)
mypi = array(0, dtype=float)

if myrank == 0:
    _n = 20 # Enter the number of intervals
    n.fill(_n)
comm.Bcast([n, MPI.INT], root=0)
_mypi = comp_pi(n, myrank, nprocs)
mypi.fill(_mypi)
comm.Reduce([mypi, MPI.DOUBLE], [pi, MPI.DOUBLE],
            op=MPI.SUM, root=0)
if myrank == 0:
    prn_pi(pi, PI)
```

## Example 3: Python with MPI: mpi4py

Submission script `sub_calc_pi.sh`:

```
#!/bin/bash
#
#SBATCH --job-name="Py_pi"
#SBATCH --time=00:10:00
#SBATCH --ntasks=4
#SBATCH --cpus-per-task=1
#SBATCH --partition=compute
#SBATCH --mem-per-cpu=1G
#SBATCH --account=Education-EEMCS-Courses-WI4049TU

module load 2023r1
module load openmpi
module load python
module load py-numpy
module load py-mpi4py

srun python calculate_pi.py
```

## Example 3: Python with MPI: mpi4py

Let's submit the job:

```
user@login01:~ $ sbatch sub_calc_pi.sh
```

And check out the result:

```
pi is approximately 3.1418009868930934, error is  
0.0002083333033003
```

## Example 4: ASE molecules generator

Let's install a new python module, called ASE:

```
#!/bin/bash

# Install ASE:
module load 2023r1
module load python
module load py-pip
module load py-numpy
module load py-scipy
module load py-matplotlib

python -m pip install --user ase
```

```
[NetID@login02 ~]$ chmod +x install_ase.sh
```

```
[NetID@login02 ~]$ ./install_ase.sh
```

## Example 4: ASE molecules generator

Now we are able to run our script:

```
# Import modules:
import os
import sys
import subprocess
import ase
from ase.io import read,write
from ase.build import molecule
from ase.optimize import BFGS
from ase.calculators.emt import EMT

def main():
    # Print info:
    info = """
    This script does the following:
    1. Creates a parent directory "molecules".
    2. Reads an array of molecules and does the following steps for each molecule:
        3. Creates a sub-directory for current molecule.
        4. Generates .xyz file containing Cartesian coordinates of current molecule.
        5. Decides if to write or to append the .log file.
        6. Optimizes the initial geometry of current molecule with EMT and ...
        6. ... writes optimization .log and ASE trajectory files for current molecule.
    """
    print(info)
```

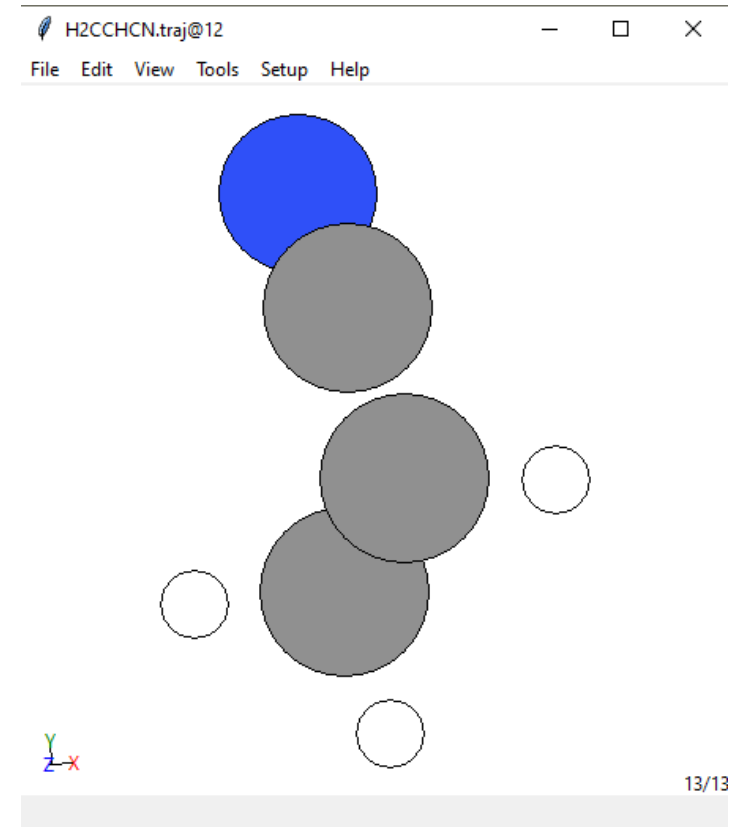
## Example 4: ASE molecules generator

```
[NetID@login02 ~]$ sbatch sub_to_queue.sh
```

```
[NetID@login02 ~]$ cat slurm-XXX.out
```

### Exercise:

Install ASE. Submit the job. Inspect the output file and generated folders.





## (Bonus) Example 5: PyTorch

Let's prepare a little program to check how many GPUs are available to us:

```
import torch
cuda_avail = torch.cuda.is_available()
if cuda_avail:
    print("Torch CUDA is available")
    num_of_devices = torch.cuda.device_count()
    if num_of_devices:
        print("Number of CUDA devices: {}".format(num_of_devices))
        current_device = torch.cuda.current_device()
        current_device_id = torch.cuda.device(current_device)
        current_device_name = torch.cuda.get_device_name(current_device)
        print("Current device id: {}".format(current_device_id))
        print("Current device name: {}".format(current_device_name))
    else:
        print("No CUDA devices!")
else:
    print("Torch CUDA is not available!")
```

## (Bonus) Example 5: PyTorch

Submission script `sub_pytorch_gpus.sh`:

```
#!/bin/bash
#
#SBATCH --job-name="PyTorch"
#SBATCH --time=00:10:00
#SBATCH --ntasks=1
#SBATCH --cpus-per-task=1
#SBATCH --gpus-per-task=1
#SBATCH --partition=gpu
#SBATCH --mem=4G
#SBATCH --account=Education-EEMCS-Courses-WI4049TU

module load 2023r1
module load openmpi
module load py-torch

srun python test_pytorch_gpus.py
```

## (Bonus) Example 5: PyTorch

Let's submit the job and check out the result:

```
Torch CUDA is available  
Number of CUDA devices: 1  
Current device id: <torch.cuda.device object at 0x155555379dc0>  
Current device name: Tesla V100S-PCIE-32GB
```

Discussion and questions

Thanks for your attention