

DELFT UNIVERSITY OF TECHNOLOGY

INTRODUCTION TO HIGH PERFORMANCE COMPUTING
WI4049TU

Lab Report

Author:
Elias Wachmann (6300421)

November 22, 2024



General Remarks

This final Lab report includes the answers for the exercises (base grad denoted in paranthesis):

0. Introductory exercise (0.5)
1. Poisson solver (1.75)
2. Finite elements simulation (1.0)
3. Eigenvalue solution by Power Method on GPU (1.75)

The optional **shining points** (e.g., performance analysis, optimization, discussion, and clarifying figures) which yield further points are usually marked by a small blue heading in the text or an additional note is added under a figure or table. For example:

This is a shining point.

0 Introductory exercise

In the introductory lab session, we are taking a look at some basic features of MPI. We start out very simple with a hello world program on two nodes.

Hello World

```
1 #include "mpi.h"
2 #include <stdio.h>
3
4 int np, rank;
5
6 int main(int argc, char **argv)
7 {
8     MPI_Init(&argc, &argv);
9     MPI_Comm_size(MPI_COMM_WORLD, &np);
10    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
11
12    printf("Node %d of %d says: Hello world!\n", rank, np);
13
14    MPI_Finalize();
15    return 0;
16 }
```

This program can be compiled with the following command:

```
mpicc -o helloworld1.out helloworld1.c
```

And run with:

```
srunc -n 2 -c 4 --mem-per-cpu=1GB ./helloworld1.out
```

We get the following output:

```
Node 0 of 2 says: Hello world!
Node 1 of 2 says: Hello world!
```

From now on I'll skip the compilation and only mention on how many nodes the program is run and what the output is / interpretation of the output.

0.a) Ping Pong

I used the template to check how long `MPI_Send` and `MPI_Recv` take. The code can be found in the appendix for this section.

I've modified the printing a bit to make it easier to gather the information. Then I piped the program output into a textfile for further processing in python. I ran it first on one and then on two nodes as specified in the

assignment sheet. Opposed to the averaging over 5 send / receive pairs, I've done 1000 pairs. Furthermore I reran the whole program 5 times to gather more data. All this data is shown in the following graph:



Figure 1: Ping Pong: Number of bytes sent vs. average time taken from 1000 pairs of send / receive. 5 runs shown for each size as scatter plot. Mean of these 5 runs shown as line. Blue small fit includes all data points up to 131072 bytes, blue large from there. Red small fit includes all data points up to 32768 bytes, red large from there.

As can be seen in the data and the fits, there are outliers especially for the larger data sizes. For our runs we get the following fits and R^2 values:

Run Type	Data Size	Fit Equation	R^2 Value
Single Node	Small (≤ 131072)	$5.95 \times 10^{-7} \cdot x + 7.97 \times 10^{-4}$	0.92
Single Node	Large (≥ 131072)	$4.61 \times 10^{-7} \cdot x + 1.23 \times 10^{-2}$	0.89
Two Node	Small (≤ 32768)	$1.07 \times 10^{-6} \cdot x + 2.60 \times 10^{-3}$	0.97
Two Node	Large (≥ 32768)	$4.41 \times 10^{-7} \cdot x + 3.42 \times 10^{-3}$	0.97

Table 1: Fit Equations and R^2 Values for Single Node and Two Node Runs

Note: Each run was performed 5 times (for 1 and 2 nodes) to get a fit on the data and calculate a R^2 value.

TODO: Further analysis needed?

Extra: Ping Pong with MPI_SendRecv

We do the same analysis for the changed program utilizing `MPI_SendRecv`. The code can be found in the appendix for this section.

We get the following graph from the measurements which were performed in the same way as for the previous program:

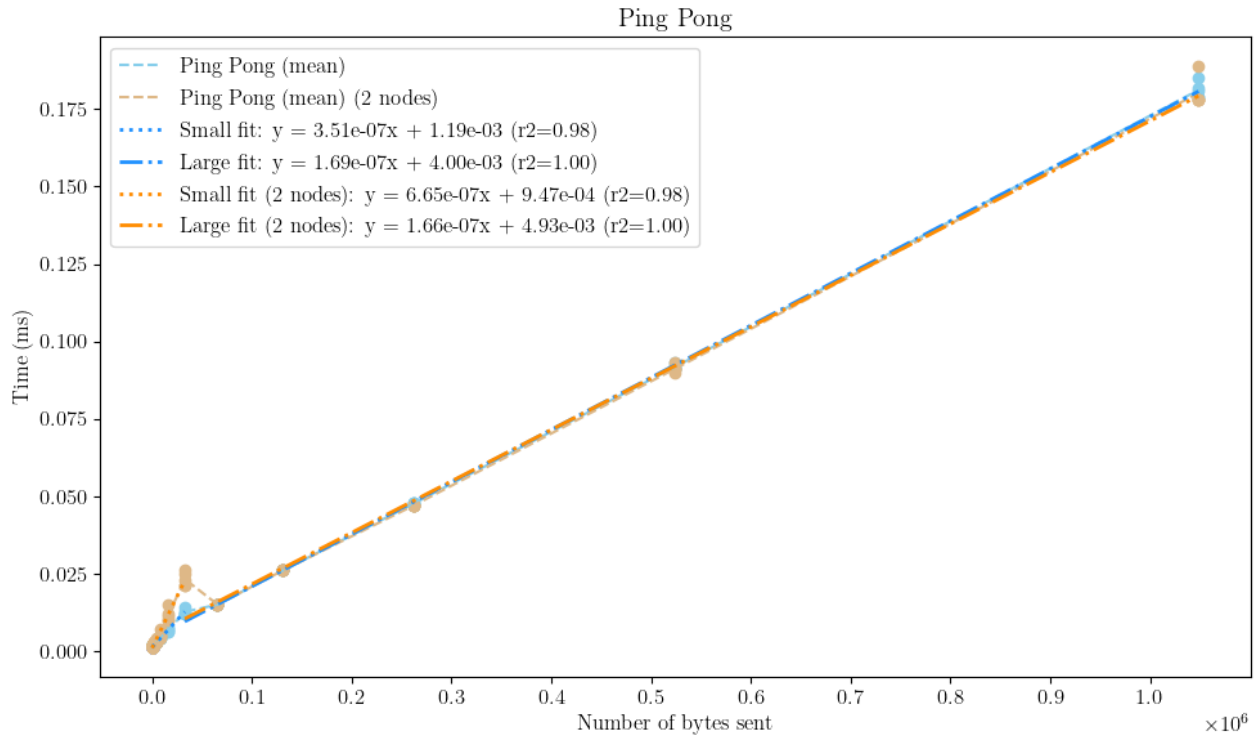


Figure 2: Ping Pong with MPI_SendRecv: Number of bytes sent vs. average time taken from 1000 pairs of send / receive. 5 runs shown for each size as scatter plot. Mean of these 5 runs shown as line. Blue small fit includes all data points up to 32768 bytes, blue large from there. Red small fit includes all data points up to 32768 bytes, red large from there.

We get the following fits and Rš values for the runs:

Run Type	Data Size	Fit Equation	Rš Value
Single Node	Small (≤ 32768)	$3.51 \times 10^{-7} \cdot x + 1.19 \times 10^{-3}$	0.98
Single Node	Large (≥ 32768)	$1.69 \times 10^{-7} \cdot x + 4.00 \times 10^{-3}$	1.00
Two Node	Small (≤ 32768)	$6.65 \times 10^{-7} \cdot x + 9.47 \times 10^{-4}$	0.98
Two Node	Large (≥ 32768)	$1.66 \times 10^{-7} \cdot x + 4.93 \times 10^{-3}$	1.00

Table 2: Fit Equations and Rš Values for Single Node and Two Node Runs

TODO: Further analysis needed?

0.b) MM-product

After an introduction of the matrix-matrix multiplication code in the next section, the measured speedups are discussed in the subsequent section.

Explanation of the code

For this exercise I've used the template provided in the assignment sheet as a base to develop my parallel implementation for a matrix-matrix multiplication. The code can be found in the appendix for this section.

The program can be run either in sequential (default) or parallel mode (parallel as a command line argument). For the sequential version, the code is practically unchanged and just refactored into a function for timing purposes. The parallel version is more complex and works as explained below:

First, rank 0 computes a sequential reference solution. Then rank 0 distributes the matrices in the following way in `splitwork`:

- Matrix A is split row-wise by dividing the number of rows by the number of nodes.
- The first worker (=rank 1) gets the most rows starting from row 0:
 $\text{total_rows} - (\text{nr_workers} - 1) \cdot \text{floor}(\frac{\text{total_rows}}{\text{nr_workers}})$.
- All other workers and the master (= rank 0) get the same number of rows: $\text{floor}(\frac{\text{total_rows}}{\text{nr_workers}})$.
- The master copies the corresponding rows of matrix A and the whole transposed matrix B* into a buffer (for details on MM_input buffer see below) for each worker and sends them off using MPI_Isend.
- The workers receive the data using MPI_Recv and then compute their part of the matrix product and send only the rows of the result matrix back to the master using MPI_Send.
- In the meanwhile the master computes its part of the matrix product.
- Using MPI_Waitall the master waits for all data to be sent to the workers and only afterwards calls MPI_Recv to gather the results from the workers.
- Finally all results are gathered by the master in the result matrix.

Assume we have a 5x5 matrix A and 2 workers (rank 1 and rank 2) and master (rank 0). The partitioning is done row-wise as follows:

Partitioning Example

$$A = \begin{pmatrix} a_{11} & a_{12} & a_{13} & a_{14} & a_{15} \\ a_{21} & a_{22} & a_{23} & a_{24} & a_{25} \\ a_{31} & a_{32} & a_{33} & a_{34} & a_{35} \\ a_{41} & a_{42} & a_{43} & a_{44} & a_{45} \\ a_{51} & a_{52} & a_{53} & a_{54} & a_{55} \end{pmatrix} \rightarrow \begin{pmatrix} \text{Worker 1} \\ \text{Worker 1} \\ \text{Worker 1} \\ \text{Master} \\ \text{Master} \end{pmatrix}$$

- **Rank 0 (Master):** Rows 4 and 5 (last two rows)
- **Rank 1 (Worker 1):** Rows 1 to 3 (first three rows) - Worker 1 always gets the most rows

This partitioning can be visually represented as:

$$\text{Master (rank 0): } \begin{pmatrix} a_{41} & a_{42} & a_{43} & a_{44} & a_{45} \\ a_{51} & a_{52} & a_{53} & a_{54} & a_{55} \end{pmatrix}$$

$$\text{Worker 1 (rank 1): } \begin{pmatrix} a_{11} & a_{12} & a_{13} & a_{14} & a_{15} \\ a_{21} & a_{22} & a_{23} & a_{24} & a_{25} \\ a_{31} & a_{32} & a_{33} & a_{34} & a_{35} \end{pmatrix}$$

Each worker computes its part of the matrix product, and the master gathers the results at the end and compiles them into the final matrix.

The MM_input buffer is used to store the rows of matrix A and the whole matrix B for each worker. It is implemented using a simple struct:

```
1 typedef struct MM_input {
2     size_t rows;
3     double *a;
4     double *b;
5 } MM_input;
```

***[Optimization] Note on transposed matrix B:** It is usually beneficial from a cache perspective to index arrays sequentially or in a row-major order. However, in the matrix-matrix multiplication, we access the elements of matrix B in a column-wise order. This leads to cache misses and is not optimal. To mitigate this, we can transpose matrix B and then access it in a row-wise order. This is done in the code by the master before sending the data to the workers.

Discussion of the speedups

The code was run on Delft's cluster with 1, 2, 4, 8, 16, 24, 32, 48, and 64 nodes. For the experiments the matrix size of A and B was set to 2000×2000 . This means that the program has to evaluate 2000 multiplications and 1999 additions for each element of the resulting matrix C . In total this results in $\approx 2000^3 = 8 \times 10^9$ operations. The command looked similar to the following for the different node counts:

```
srun -n 48 --mem-per-cpu=4GB --time=00:02:00 ./MM.out parallel
```

For this experiment, the execution time was measured and the speedup was calculated. The results are shown in [Table 3](#) and [Figure 3](#).

CPU Count	Execution Time / s	Approx. Speedup
1	47.11	1.0
2	10.26	4.6
4	10.30	4.6
8	5.20	9.1
16	2.97	15.9
24	2.54	18.5
32	2.29	20.6
48	2.98	15.8
64	1.72	27.4

Table 3: Execution Time vs CPU Count



Figure 3: Speedup vs CPU Count
Black \times marks the average of the rerun for $n = 48$.

Note: The speedup is calculated as $S = \frac{T_1}{T_p}$, where T_1 is the execution time on 1 node and T_p is the execution time on p nodes.

Discussion:

As one can clearly discern from the data in [Table 3](#) and [Figure 3](#), the speedup increases with the number of nodes (with the exception of $n = 48$). This is expected as the more nodes we have, the more work can be done in

parallel. However, the speedup is not linear. This is due to the overhead of communication between the nodes. The more nodes we have, the more communication is needed, and this overhead increases. This is especially visible in the data for $n = 48$. Here the speedup is lower than for $n = 32$. For this run the communication didn't went as smooth as for the other runs. This can potentially be attributed to the fact that one (or more) of the nodes or the network was under heavy load during this task.

[Further investigation] After observing this slower speed for the $n = 48$, I reran the tests multiple times and got a runtime of around 1.9s which was to be expected initially. Therefore, this one run is an odd one out, most likely due to the reasons mentioned above! I've also added the averaged data of the reruns as a datapoint in [Figure 3](#).

Another interesting fact can be seen when comparing the time taken for $n = 1$ and $n = 2$. They don't at all scale with the expected factor of 2. This is could be due to the fact, that the resource management system prefers runs with multiple nodes instead of a single node (= sequential).

Additional notes: The flag `-mem-per-cpu=<#>GB` was set depending on the number of nodes used. For 1-24 nodes 8GB was used, for 32-48 nodes 4GB, and for 64 nodes 3GB. This had to be done to comply with QOS policy on the cluster.

TODO: Data locality?

1 Poisson solver

In this section of the lab report, we will discuss a parallel implementation of the Poisson solver. The Poisson solver is a numerical method used to solve the Poisson equation, which is a partial differential equation that is useful in many areas of physics.

Note: For local testing and development I'll run the code with `mpirun` instead of the `srun` command on the cluster.

1.1 Building a parallel Poisson solver

For the first part of the exercise we follow the steps lined out in the assignment sheet. I'll comment on the steps 1 through 10 and related questions below. The finished implementation can be found in the appendix for this section.

1. **Step:** After adding `MPI_Init` and `MPI_Finalize`, we can run the program with multiple processes. We can see that the program runs with 4 processes in [Figure 4](#) via the quadrupled output.

```
etschgi1@Deep-Thought:~/REPOS/HPC/01_lab1/src$ mpirun -np 4 ./mpi.out
Number of iterations : 2355
Number of iterations : 2355
Number of iterations : 2355
Elapsed procestime:      0.133189 s
Number of iterations : 2355
Elapsed procestime:      0.134150 s
Elapsed procestime:      0.134474 s
Elapsed procestime:      0.135356 s
```

Figure 4: MPI_Poisson after Step 1 - Running with 4 processes

2. **Step:** To see which process is doing what, I included the rank of the process for the print statements as shown in [Figure 5](#).

```
etschgi1@Deep-Thought:~/REPOS/HPC/02_lab1/src$ mpirun -np 4 ./mpi.out
(0) Number of iterations : 2355
(2) Number of iterations : 2355
(0) Elapsed procestime:      0.135963 s
(2) Elapsed procestime:      0.137101 s
(3) Number of iterations : 2355
(3) Elapsed procestime:      0.139614 s
(1) Number of iterations : 2355
(1) Elapsed procestime:      0.142026 s
```

Figure 5: MPI_Poisson after Step 2 - Running with 4 processes

3. **Step:** Next we define `wtime` as a global double and replace the four utility timing functions with the ones given on Brightspace. A quick verification as shown in [Figure 6](#) shows that the program still runs as expected.

```
etschgi1@Deep-Thought:~/REPOS/HPC/02_lab1/src$ mpirun -np 4 ./mpi.out
(3) Number of iterations : 2355
(1) Number of iterations : 2355
(3) Elapsed Wtime      0.134918 s ( 98.5% CPU)
(1) Elapsed Wtime      0.134459 s ( 98.5% CPU)
(0) Number of iterations : 2355
(2) Number of iterations : 2355
(0) Elapsed Wtime      0.138669 s ( 98.5% CPU)
(2) Elapsed Wtime      0.138910 s ( 98.5% CPU)
```

Figure 6: MPI_Poisson after Step 3 - Running with 4 processes

4. **Step:** Next we check if two processes indeed give the same output. Both need 2355 iterations to converge and the `diff` command returned no output, which means that the files content is identical.
5. **Step:** Now only the process with rank 0 will read data from files and subsequently broadcast it to the others. Testing this again with 2 processes, we see an empty diff of the output files and the same number of iterations needed to converge.

6. **Step:** We create a cartesian grid of processes using `MPI_Cart_create` and use `MPI_Cart_shift` to find the neighbors of each process. We can see that the neighbors are correctly identified in [Figure 7](#).

```
(0) (x,y)=(0,0)
(0) top 1, right -2, bottom -2, left 2
(1) (x,y)=(0,1)
(1) top -2, right -2, bottom 0, left 3
(2) (x,y)=(1,0)
(2) top 3, right 0, bottom -2, left -2
(3) (x,y)=(1,1)
(3) top -2, right 1, bottom 2, left -2
```

Figure 7: MPI_Poisson after Step 6 - Running with 4 processes on a 2x2 grid

When there is no neighbor in a certain direction, -2 (or `MPI_PROC_NULL`) is returned.

7. **Step:** We overhaul the setup to get a proper local grid for each process. Furthermore, we only save the relevant source fields in the local grid for each process.

With for instance 3 processes you should see that 1 or 2 processes do not do any iteration. Do you understand why?

If we have a look at the input file we see that there are only 3 source fields in the grid. This means that the process that does not have a source field in its local grid will not do any iterations (or only 1). Therefore, if we have 3 processes and the distribution of source fields as given in the input file only 1 process will do iterations if processes are ordered in x-direction and 2 if ordered in y-direction. From this we can conclude that indeed all processes have different local grids and perform different calculations.

```
etschgi1@Deep-Thought:~/REPOS/HPC/02_lab1/src$ mpirun -np 3 ./mpi.out 3 1
(0) (x,y)=(0,0)
(0) top -2, right -2, bottom -2, left 1
(1) (x,y)=(1,0)
(1) top -2, right 0, bottom -2, left 2
(2) (x,y)=(2,0)
(2) top -2, right 1, bottom -2, left -2
(0) Number of iterations : 1
(2) Number of iterations : 1
(2) Elapsed Wtime 0.000668 s ( 95.3% CPU)
(0) Elapsed Wtime 0.000917 s ( 95.9% CPU)
(1) Number of iterations : 695
(1) Elapsed Wtime 0.014772 s ( 95.2% CPU)
```

```
etschgi1@Deep-Thought:~/REPOS/HPC/02_lab1/src$ mpirun -np 3 ./mpi.out 1 3
(1) (x,y)=(0,1)
(1) top 2, right -2, bottom 0, left -2
(1) Number of iterations : 1
(2) (x,y)=(0,2)
(2) top -2, right -2, bottom 1, left -2
(0) (x,y)=(0,0)
(0) top 1, right -2, bottom -2, left -2
(1) Elapsed Wtime 0.000616 s ( 95.4% CPU)
(0) Number of iterations : 601
(2) Number of iterations : 723
(0) Elapsed Wtime 0.017636 s ( 95.3% CPU)
(2) Elapsed Wtime 0.017801 s ( 95.3% CPU)
```

Figure 8: MPI_Poisson after Step 7 - Running with 3 processes on a 3x1 (left) vs. 1x3 (right) grid
For the 3x1 grid, only rank 1 does iterations (> 1), for the 1x3 grid, ranks 0 and 2 do iterations (> 1).

8. **Step:** After defining and committing two special datatypes for vertical and horizontal communication, we setup the communication logic to exchange the boundary values between the processes. We call our `Exchange_Borders` function after each iteration (for both red / black grid points). Now we face the problem in which some processes may stop instantly (no source in their local grid). They will not supply any data to their neighbors, which will cause the program to hang. We shall fix this in the next step.
9. **Step:** Finally we need to implement the logic to check for convergence (in a global sense). We do this by using a `MPI_Allreduce` call with the `MPI_MAX` operation. This way we aggregate all deltas and choose the biggest one for the global delta which we use in the while-loop-condition to check for convergence. We can see that the program now runs as expected in [Figure 9](#).

```

(0) (x,y)=(0,0)
(0) top -1, right 2, bottom 1, left -1
(1) (x,y)=(0,1)
(1) top 0, right 3, bottom -1, left -1
(2) (x,y)=(1,0)
(2) top -1, right -1, bottom 3, left 0
(3) (x,y)=(1,1)
(3) top 2, right -1, bottom -1, left 1
(0) Number of iterations : 2355
(1) Number of iterations : 2355
(2) Number of iterations : 2355
(3) Number of iterations : 2355
(1) Elapsed Wtime      0.287549 s ( 99.9% CPU)
(2) Elapsed Wtime      0.287537 s (100.0% CPU)
(3) Elapsed Wtime      0.287537 s (100.0% CPU)
(0) Elapsed Wtime      0.295957 s ( 99.9% CPU)

```

Figure 9: MPI_Poisson after Step 9 - Running with 4 processes on a 2x2 grid

Note that this run in Figure 9 was done with another pc and another MPI implementation. Therefore, we see -1 for cells without a neighbor! However, other than that cosmetic difference it has no impact on the programm.

10. **Step:** Now we only have to fix two remaining things. First we have to make sure that each process uses the right global coordinates for the output file in the end. Therefore, we change the function a bit to include the specific x-/y-offset for each processor. The second thing is the potential problem, that different processors might start with different (red/black) parities. In order to accomplish a global parity we simply have to change the calculation in the if in Do_Step from

```

1  if ((x + y) % 2 == parity && source[x][y] != 1)

```

to

```

1  if ((x + offset[X_DIR] + y + offset[Y_DIR]) % 2 == parity && source[x][y] != 1)

```

this guarantees that during a given iteration all processors are using the same parity.

This just leaves one question open: Are the results acutally the same?

Checking the output files of the MPI-implementation with the sequential reference indeed shows identical numerical values for the calculated points. Furthermore, the needed iterationcount is also identical which isn't a big surprise, given that the two programmes perform the exact same calculation steps.

1.2 Exercises, modifications, and performance aspects

For this subsection we'll define the following shorthand notation:

n :	the number of iterations
g :	gridsize
t :	time needed in seconds
pt :	processor topology in form pxy , where:
p :	number of processors used
x :	number of processors in x-direction
y :	number of processors in y-direction

Table 4: Notation for this section

$pt = 414$ means 4 processors in a 1×4 topology.

Note on different Versions:

For the following exercises the implementation will be slightly adapted to measure different performance aspects. To facilitate this, we will use defines to switch between different versions of the code at compile time. The final version of the poissonsolver can be found in the appendix for this section.

1.2.1 Over-relaxation (SOR)

We start off by rewriting the `Do_Step` routine to facilitate SOR updates. Furthermore, we need h^2 , the grid spacing (which is 1 in our case) and the relaxation parameter ω to calculate the updated values. A quick test shows a speedup of roughly a factor of 10. More systematic tests will be done in the next section.

1.2.2 Optimal ω for 4 proc. on a 4x1 grid

With the power of a little python scripting we can easily test different values for ω and plot the results as seen in [Figure 10](#).

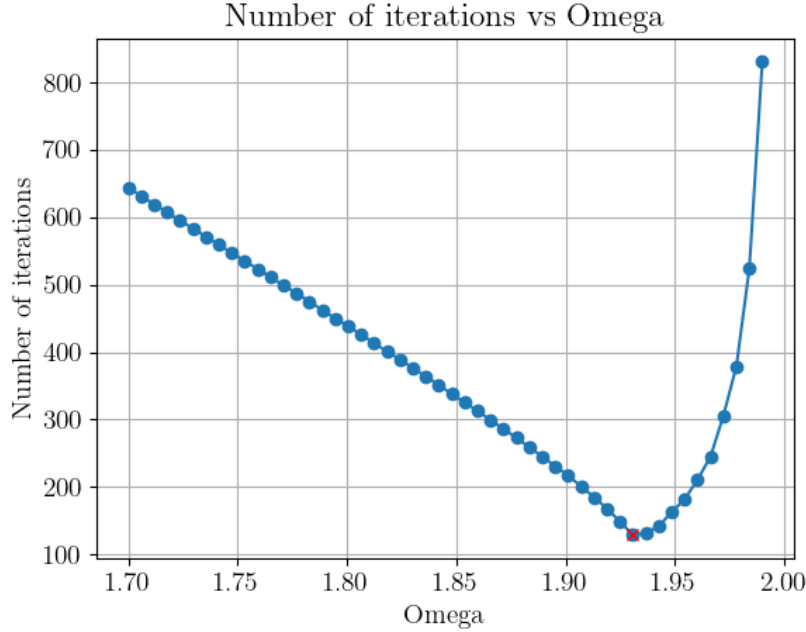


Figure 10: **Optimal ω for 4 processors on a 4x1 grid**

We find that the optimal ω is at about 1.93 for this setup with only 129 iterations. This constitutes a speedup of about 1825% compared to the sequential implementation.

N.B.: If not stated otherwise, we will use $\omega = 1.93$ for the following exercises.

1.2.3 Scaling behavior with larger grids

This investigation is carried out twice: Once with a 4×1 topology (as in the previous section) and once with a 2×2 topology. We use grid sizes of 10×10 , 25×25 , 50×50 , 100×100 , 200×200 , 400×400 , 800×800 and 1600×1600 and set $\omega = 1.95$ for all runs. The results are shown in [Figure 11](#).

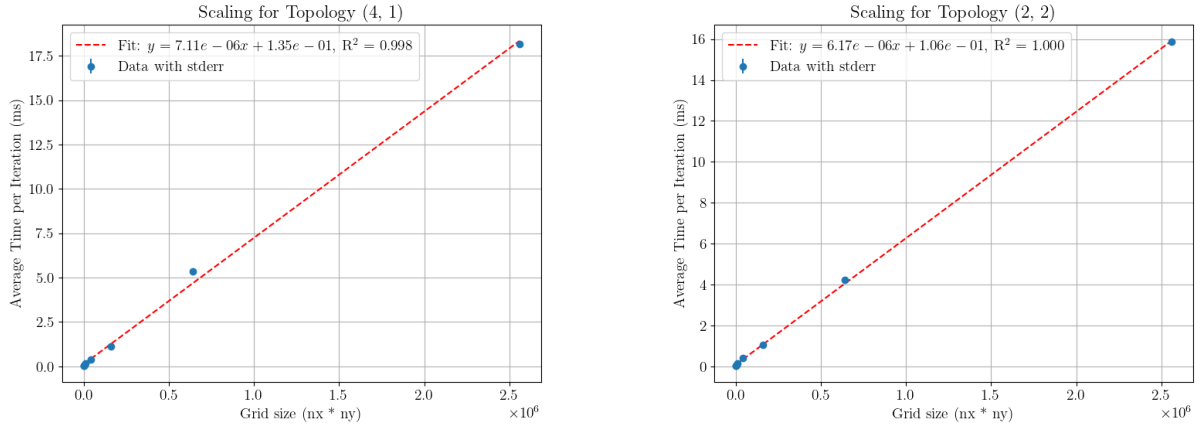


Figure 11: Scaling behavior of the Poisson solver with different grid sizes and processor topologies

As seen by the high R^2 values in the plots, the scaling behavior is very close to linear. We obtain the following scaling factors for the different grid sizes and topologies from the linear fits:

Topology	α	β
4×1	$1.35 + 10^{-1}$	$7.11 + 10^{-6}$
2×2	$1.06 + 10^{-1}$	$6.17 + 10^{-6}$

Table 5: Scaling factors for different processor topologies for the Poisson solver

Using: $t(n) = \alpha + \beta \cdot n$ as a model

What can you conclude from the scaling behavior?

We see that the scaling behavior is very close to linear for both topologies. This means that the parallel implementation scales as expected with the number of grid points.

If we compare the scaling factors (β) for the two topologies we see that the 2×2 topology scales slightly better than the 4×1 topology. This is not surprising, as the 2×2 topology has a more balanced communication workload balance. In the 2×2 topology every processor has two neighbors, while in the 4×1 topology the processors at the ends only have one neighbor. This is a general trend: A topology which divides the grid into square / square-like parts will scale better than a topology which divides the grid into long and thin parts.

In essence: We want to keep the communication between processors as balanced as possible to achieve the best scaling behavior.

1.2.4 Scaling behavior [Theory - no measurements]

If I could choose between a 16×1 , 8×2 , 4×4 , 2×8 , 1×16 topology, I would choose the 4×4 topology. This is because the 4×4 topology has the most balanced communication workload balance, as detailed in the [Shining](#) in [subsubsection 1.2.3](#).

1.2.5 Iterations needed for convergence scaling

We investigate the number of iterations needed for convergence using the 4×1 topology square grids with sidelength: 10, 25, 50, 100, 200, 400, 800, 1600. The results for different ω are shown in [Figure 12](#).

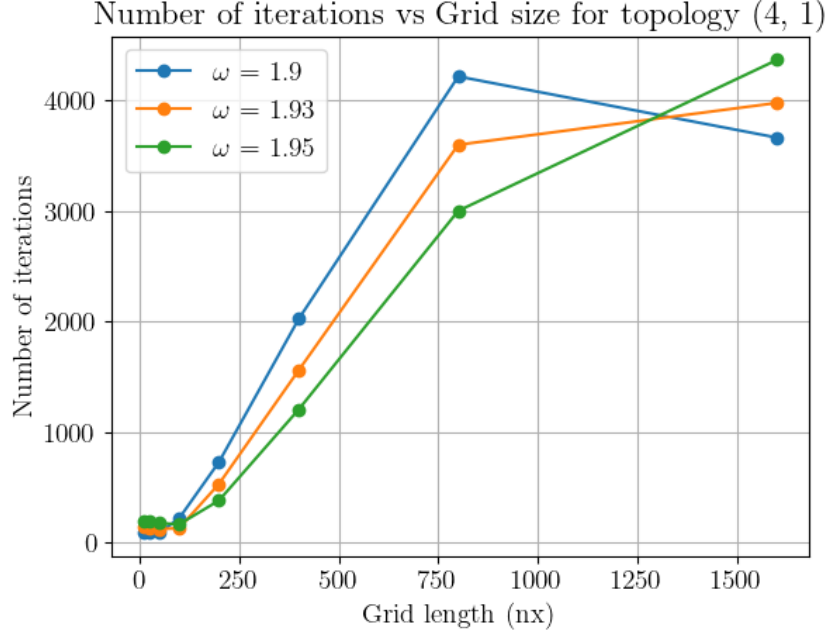


Figure 12: Iterations needed for convergence with different grid side lengths

We can clearly see that the number of iterations till convergence increases with the problem size. At first, I expected linear growth proportional to the number of gridpoints. However, it turns out that the number of iterations actually grow slower and in a square root like fashion. This can be seen by the linear behavior in the plot of grid-side length against iterations.

Why is the number of iterations needed for convergence $\propto \sqrt{g}$?

Our poisson problem is a discretized system in 2D space. The condition number of the matrix we have to solve is proportional to the number of gridpoints in our system. SOR uses the spectral properties of the matrix to solve in a way such that the dominant error mode takes time proportional to the diameter of the domain to converge. This means it is proportional to $\sqrt{g} = \sqrt{n_x \cdot n_y}$.

Why does omega with the best performance change with the grid size?

As can be seen in Figure 12 $\omega = 1.9$ beats the other two values for very small and the largest gridsize. For different gridsizes we get differently sized matrices we have to solve. SOR overrelaxes high-frequency errors and underrelaxes low-frequency errors (the later for stability). The optimal ω is therefore dependent on the gridsize and the error modes present in the system. In our current example, it might be that $\omega = 1.9$ is a good compromise for the grid sizes we are looking at.

1.2.6 Error as a function of the iteration number

With the same 4×1 topology and grid sizes of 800×800 the error for 15000 iterations is tracked using $\omega = 1.93$. The results are shown in Figure 13.

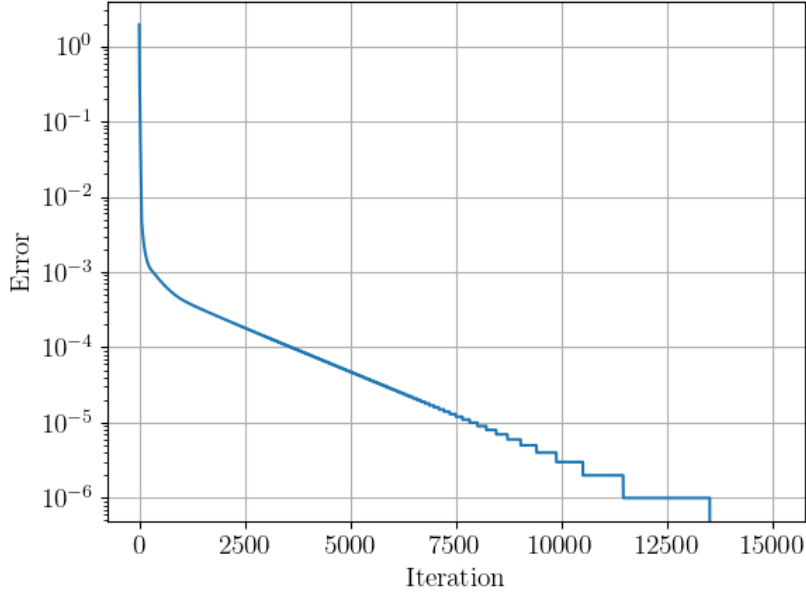


Figure 13: Error as a function of the iteration number

At first the error decreases rapidly in the first few iterations to about 10^{-3} (logarithmic scale!). After that the error decreases more slowly until it is below floating point precision. **Note:** All calculations are done using double precision floating point numbers and only the error recording was done using single precision which leaves the step-like artifacts in the plot.

1.2.7 Optional - Gain performance by reducing MPI_Allreduce calls

The last subsection showed us that the error reduces monotonically. We might be able to save some time by leaving out some checks and maybe check the global error every 10th or 100th iteration only.

First, we should benchmark if it is at all wise to optimize here, by measuring how long the MPI_Allreduce call takes. We can do this by measuring the time needed for the MPI_Allreduce call in the Do_Step function and summing up to get the total time spent in MPI_Allreduce calls.

We again solve with a 4×1 topology, $\omega = 1.93$ and a 800×800 grid: It takes roughly 20 seconds of which the processors spend around 1 - 2 seconds in the MPI_Allreduce call. This is a significant amount of time ($(7.0 \pm 0.4)\%$). This means we would save some time by reducing the number of MPI_Allreduce calls and calculating 9 (0.25% of total) more iterations wouldn't hurt us too bad because it takes 3601 to converge!

We run the program three times with MPI_Allreduce calls every 1, 10 and 100 iterations and get the speedups in MPI_Allreduce calls as shown in Table 6.

Iterations	MPI_Allreduce - speedup (factor)	calculated overall speedup (%)
1	1.00	0
10	6.0 ± 2.0	5.9 ± 0.5
100	62 ± 6	6.9 ± 0.4

Table 6: Speedup in MPI_Allreduce calls for different iteration counts and calculated overall speedup (%)

As can be clearly seen from the table we can gain around 6 % using MPI_Allreduce calls every 10 iterations and around 7 % using MPI_Allreduce calls every 100 iterations. This is a significant speedup for a very small change in the code.

Note: The speedup is calculated to account for fluctuations in the runtime of the program, due to other processes running on the same machine / cluster.

1.2.8 Reduce border communication

Another way to reduce communication overhead is to reduce the number of border exchanges. To investigate if this yields a speedup we run the program on a 4×1 topology, $\omega = 1.93$ and different grid sizes and track the iterations and time as seen in Figure 14.

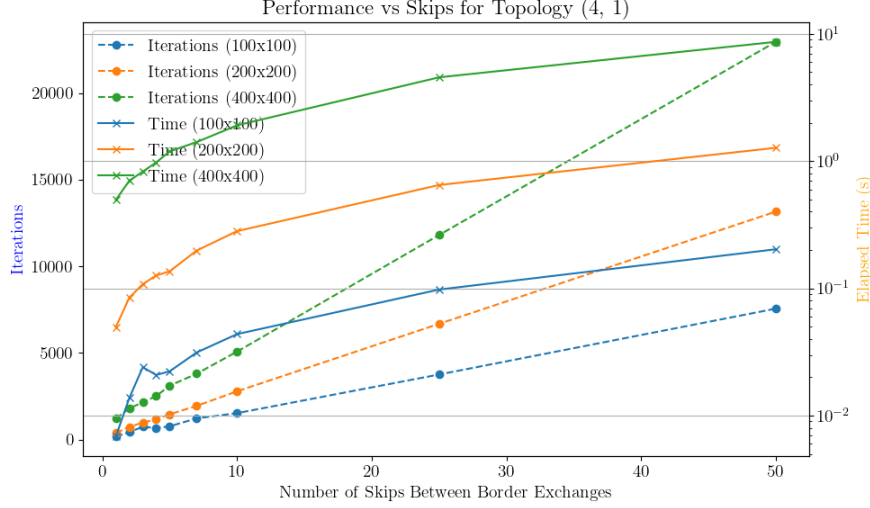


Figure 14: Speedup by reducing border exchanges

Running the with different numbers of skipped border exchanges naturally slows down convergence, meaning we need more iterations to reach the same error. For all tested grid sizes the initial SOR version without skipping border exchanges has the fewest iterations needed to convergence and also the fastest runtime.

What can you conclude from the results?

We can conclude that reducing the number of border exchanges does not yield a speedup. The reason for this is that we have to calculate more iterations to converge to the solution which outweighs the gains from reduced communication overhead. Interestingly, for the 100×100 grid there exists a local minimum in time at 4 skipped border exchanges compared to 3 skipped. This is likely due to our source field distribution and thus specific to our problem.

1.2.9 Optimize Do_Step loop

In `Do_Step` we iterate over the whole grid but only update one of the two parities at a time. This means we can split the loop into two loops, one for each parity. We start out with something like this:

```

1  for (x = 1; x < dim[X_DIR] - 1; x++){
2      for (y = 1; y < dim[Y_DIR] - 1; y++){
3          if ((x + offset[X_DIR] + y + offset[Y_DIR]) % 2 == parity && source[x][y] != 1){
4              ...

```

and we change it to:

```

1  int start_y;
2  for (x = 1; x < dim[X_DIR] - 1; x++){
3      start_y = ((1 + x + offset[X_DIR] + offset[Y_DIR]) % 2 == parity) ? 1 : 2;
4      for (y = start_y; y < dim[Y_DIR] - 1; y += 2){
5          if (source[x][y] != 1){
6              ...

```

The basic idea is to avoid y-coordinates which are not in the parity we are currently updating. We measure 10 runs for a 800×800 grid and a 4×1 topology with $\omega = 1.93$ and get the following times:

$$t_{\text{no_improvements}} = (5.59 \pm 0.05) \text{ s} \quad \text{and} \quad t_{\text{loop_improvements}} = (4.64 \pm 0.07) \text{ s}$$

So we get a minimal speedup of about 17% by optimizing the loop which is a enormous speedup for such a small change.

Why does this make such a difference

The reason for this is that we avoid unnecessary looping and if statements. This means that we have less overhead in the loop and can therefore calculate faster by skipping the unnecessary loop entries.

1.2.10 Optional - Time spent within `Exchange_Borders`

We can measure the time spent in `Exchange_Borders` by adding a timer to the function. We run the program with a 4×1 topology, $\omega = 1.93$ and a 800×800 grid and get the following times:

2 Finite elements simulation

3 Eigenvalue solution by Power Method on GPU

Appendix - Introductory exercise

The following code was used for the ping pong task:

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <mpi.h>
4
5 // Maximum array size 2^20= 1048576 elements
6 #define MAX_EXPONENT 20
7 #define MAX_ARRAY_SIZE (1<<MAX_EXPONENT)
8 #define SAMPLE_COUNT 1000
9
10 int main(int argc, char **argv)
11 {
12     // Variables for the process rank and number of processes
13     int myRank, numProcs, i;
14     MPI_Status status;
15
16     // Initialize MPI, find out MPI communicator size and process rank
17     MPI_Init(&argc, &argv);
18     MPI_Comm_size(MPI_COMM_WORLD, &numProcs);
19     MPI_Comm_rank(MPI_COMM_WORLD, &myRank);
20
21
22     int *myArray = (int *)malloc(sizeof(int)*MAX_ARRAY_SIZE);
23     if (myArray == NULL)
24     {
25         printf("Not enough memory\n");
26         exit(1);
27     }
28     // Initialize myArray
29     for (i=0; i<MAX_ARRAY_SIZE; i++)
30         myArray[i]=1;
31
32     int number_of_elements_to_send;
33     int number_of_elements_received;
34
35     // PART C
36     if (numProcs < 2)
37     {
38         printf("Error: Run the program with at least 2 MPI tasks!\n");
39         MPI_Abort(MPI_COMM_WORLD, 1);
40     }
41     double startTime, endTime;
42
43     // TODO: Use a loop to vary the message size
44     for (size_t j = 0; j <= MAX_EXPONENT; j++)
45     {
46         number_of_elements_to_send = 1<<j;
47         if (myRank == 0)
48         {
49             myArray[0]=myArray[1]+1; // activate in cache (avoids possible delay when sending
the 1st element)
50             startTime = MPI_Wtime();
51             for (i=0; i<SAMPLE_COUNT; i++)
52             {
53                 MPI_Send(myArray, number_of_elements_to_send, MPI_INT, 1, 0,
54                     MPI_COMM_WORLD);
55                 MPI_Probe(MPI_ANY_SOURCE, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
56                 MPI_Get_count(&status, MPI_INT, &number_of_elements_received);
57
58                 MPI_Recv(myArray, number_of_elements_received, MPI_INT, 1, 0,
59                     MPI_COMM_WORLD, MPI_STATUS_IGNORE);
60             } // end of for-loop
61
62             endTime = MPI_Wtime();
63             printf("Rank %2.1i: Received %i elements: Ping Pong took %f seconds\n", myRank,
number_of_elements_received, (endTime - startTime)/(2*SAMPLE_COUNT));
64         }
65         else if (myRank == 1)
66         {
67             // Probe message in order to obtain the amount of data
68             MPI_Probe(MPI_ANY_SOURCE, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
```

```

69     MPI_Get_count(&status, MPI_INT, &number_of_elements_received);
70
71     for (i=0; i<SAMPLE_COUNT; i++)
72     {
73         MPI_Recv(myArray, number_of_elements_received, MPI_INT, 0, 0,
74                 MPI_COMM_WORLD, MPI_STATUS_IGNORE);
75         MPI_Send(myArray, number_of_elements_to_send, MPI_INT, 0, 0,
76                 MPI_COMM_WORLD);
77     } // end of for-loop
78 }
79 }
80
81 // Finalize MPI
82 MPI_Finalize();
83
84 return 0;
85 }

```

For the bonus task, the following code was used:

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <mpi.h>
4
5  // Maximum array size 2^20= 1048576 elements
6  #define MAX_EXPONENT 20
7  #define MAX_ARRAY_SIZE (1<<MAX_EXPONENT)
8  #define SAMPLE_COUNT 1000
9
10 int main(int argc, char **argv)
11 {
12     // Variables for the process rank and number of processes
13     int myRank, numProcs, i;
14     MPI_Status status;
15
16     // Initialize MPI, find out MPI communicator size and process rank
17     MPI_Init(&argc, &argv);
18     MPI_Comm_size(MPI_COMM_WORLD, &numProcs);
19     MPI_Comm_rank(MPI_COMM_WORLD, &myRank);
20
21
22     int *myArray = (int *)malloc(sizeof(int)*MAX_ARRAY_SIZE);
23     if (myArray == NULL)
24     {
25         printf("Not enough memory\n");
26         exit(1);
27     }
28     // Initialize myArray
29     for (i=0; i<MAX_ARRAY_SIZE; i++)
30         myArray[i]=1;
31
32     int number_of_elements_to_send;
33     int number_of_elements_received;
34
35     // PART C
36     if (numProcs < 2)
37     {
38         printf("Error: Run the program with at least 2 MPI tasks!\n");
39         MPI_Abort(MPI_COMM_WORLD, 1);
40     }
41     double startTime, endTime;
42
43     // TODO: Use a loop to vary the message size
44     for (size_t j = 0; j <= MAX_EXPONENT; j++)
45     {
46         number_of_elements_to_send = 1<<j;
47         if (myRank == 0)
48         {
49             myArray[0]=myArray[1]+1; // activate in cache (avoids possible delay when sending
the 1st element)
50             startTime = MPI_Wtime();
51             for (i=0; i<SAMPLE_COUNT; i++)
52             {
53                 MPI_Sendrecv(myArray, number_of_elements_to_send, MPI_INT, 1,0,myArray,

```

```

    number_of_elements_to_send, MPI_INT, 1, 0, MPI_COMM_WORLD, &status);
54     }
55
56     endTime = MPI_Wtime();
57     printf("Rank %2.i: Received %i elements: Ping Pong took %f seconds\n", myRank,
    number_of_elements_to_send, (endTime - startTime)/(2*SAMPLE_COUNT));
58     }
59     else if (myRank == 1)
60     {
61         for (i=0; i<SAMPLE_COUNT; i++)
62         {
63             MPI_Sendrecv(myArray, number_of_elements_to_send, MPI_INT, 0,0,myArray,
    number_of_elements_to_send, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
64         }
65     }
66 }
67
68 // Finalize MPI
69 MPI_Finalize();
70
71 return 0;
72 }

```

The matrix multiplication used the following code:

```

1  /*****
2  * FILE: mm.c
3  * DESCRIPTION:
4  *   This program calculates the product of matrix a[nra][nca] and b[nca][ncb],
5  *   the result is stored in matrix c[nra][ncb].
6  *   The max dimension of the matrix is constraint with static array
7  *   declaration, for a larger matrix you may consider dynamic allocation of the
8  *   arrays, but it makes a parallel code much more complicated (think of
9  *   communication), so this is only optional.
10 *
11 *****/
12
13 #include <math.h>
14 #include <mpi.h>
15 #include <stdbool.h>
16 #include <stdio.h>
17 #include <stdlib.h>
18 #include <string.h>
19
20 #define NRA 2000 /* number of rows in matrix A */
21 #define NCA 2000 /* number of columns in matrix A */
22 #define NCB 2000 /* number of columns in matrix B */
23 // #define N 1000
24 #define EPS 1e-9
25 #define SIZE_OF_B NCA*NCB*sizeof(double)
26
27 bool eps_equal(double a, double b) { return fabs(a - b) < EPS; }
28
29 void print_flattened_matrix(double *matrix, size_t rows, size_t cols, int rank) {
30     printf("[%d]\n", rank);
31     for (size_t i = 0; i < rows; i++) {
32         for (size_t j = 0; j < cols; j++) {
33             printf("%10.2f ", matrix[i * cols + j]); // Accessing element in the 1D array
34         }
35         printf("\n"); // Newline after each row
36     }
37 }
38
39 int checkResult(double *truth, double *test, size_t Nr_col, size_t Nr_rows) {
40     for (size_t i = 0; i < Nr_rows; ++i) {
41         for (size_t j = 0; j < Nr_col; ++j) {
42             size_t index = i * Nr_col + j;
43             if (!eps_equal(truth[index], test[index])) {
44                 return 1;
45             }
46         }
47     }
48     return 0;
49 }

```

```

50
51 typedef struct {
52     size_t rows;
53     double *a;
54     double *b;
55 } MM_input;
56
57 char* getbuffer(MM_input *in, size_t size_of_buffer){
58     char* buffer = (char*)malloc(size_of_buffer * sizeof(char));
59     if (buffer == 0)
60     {
61         printf("Buffer couldn't be allocated.");
62         return NULL;
63     }
64     size_t offset = 0;
65     memcpy(buffer + offset, &in->rows, sizeof(size_t));
66     offset += sizeof(size_t);
67     size_t matrix_size = in->rows * NCA * sizeof(double);
68     memcpy(buffer + offset, in->a, matrix_size);
69     offset += matrix_size;
70     memcpy(buffer + offset, in->b, NCA*NCB*sizeof(double));
71     return buffer;
72 }
73
74 MM_input* readbuffer(char* buffer, size_t size_of_buffer){
75     MM_input *mm = (MM_input*)malloc(sizeof(MM_input));
76
77     mm->rows = ((size_t*)buffer)[0];
78     size_t offset = sizeof(size_t);
79     size_t matrix_size = mm->rows * NCA;
80     mm->a = (double*)malloc(sizeof(double)*matrix_size);
81     mm->b = (double*)malloc(sizeof(double)*matrix_size);
82     memcpy(mm->a, &(buffer[offset]), matrix_size);
83     offset += matrix_size;
84     memcpy(mm->b, &(buffer[offset]), NCA*NCB*sizeof(double));
85     free(buffer);
86     return mm;
87 }
88
89
90 void setupMatrices(double (*a)[NCA], double (*b)[NCB], double (*c)[NCB]){
91     for (size_t i = 0; i < NRA; i++) {
92         for (size_t j = 0; j < NCA; j++) {
93             a[i][j] = i + j;
94         }
95     }
96
97     for (size_t i = 0; i < NCA; i++) {
98         for (size_t j = 0; j < NCB; j++) {
99             b[i][j] = i * j;
100         }
101     }
102
103     for (size_t i = 0; i < NRA; i++) {
104         for (size_t j = 0; j < NCB; j++) {
105             c[i][j] = 0;
106         }
107     }
108 }
109
110 double multsum(double* a, double* b_transposed, size_t size){
111     double acc = 0;
112     for (size_t i = 0; i < size; i++)
113     {
114         acc += a[i]*b_transposed[i];
115     }
116     return acc;
117 }
118
119 double productSequential(double *res) {
120     // dynamically allocate to not run into stack overflow - usually stacks are
121     // 8192 bytes big -> 1024 doubles but we have 1 Mio. per matrix
122     double(*a)[NCA] = malloc(sizeof(double) * NRA * NCA);

```

```

123 double(*b)[NCB] = malloc(sizeof(double) * NCA * NCB);
124 double(*c)[NCB] = malloc(sizeof(double) * NRA * NCB);
125
126 /** Initialize matrices */
127 setupMatrices(a,b,c);
128
129 /* Parallelize the computation of the following matrix-matrix
130 multiplication. How to partition and distribute the initial matrices, the
131 work, and collecting final results.
132 */
133 // multiply
134 double start = MPI_Wtime();
135 for (size_t i = 0; i < NRA; i++) {
136     for (size_t j = 0; j < NCB; j++) {
137         for (size_t k = 0; k < NCA; k++) {
138             res[i * NCB + j] += a[i][k] * b[k][j];
139         }
140     }
141 }
142 /* perform time measurement. Always check the correctness of the parallel
143 results by printing a few values of c[i][j] and compare with the
144 sequential output.
145 */
146 double time = MPI_Wtime()-start;
147 free(a);
148 free(b);
149 free(c);
150 return time;
151 }
152
153 double splitwork(double* res, size_t num_workers){
154     if (num_workers == 0) // sadly noone will help me :(
155     {
156         printf("Run sequential!\n");
157         return productSequential(res);
158     }
159
160     double(*a)[NCA] = malloc(sizeof(double) * NRA * NCA);
161     double(*b)[NCB] = malloc(sizeof(double) * NCA * NCB);
162     double(*c)[NCB] = malloc(sizeof(double) * NRA * NCB);
163     // Transpose matrix b to make accessing columns easier - in row major way - better cache
164     // performance
165     setupMatrices(a,b,c);
166
167     double start_time = MPI_Wtime();
168     double (*b_transposed)[NCA] = malloc(sizeof(double) * NCA * NCB);
169     for (size_t i = 0; i < NCA; i++) {
170         for (size_t j = 0; j < NCB; j++) {
171             b_transposed[j][i] = b[i][j];
172         }
173     }
174
175     /** Initialize matrices */
176     // given number of workers I'll split
177     size_t rows_per_worker = NRA / (num_workers+1); //takes corresponding columns from other
178     // matrix
179     printf("rows per worker: %zu\n", rows_per_worker);
180     size_t row_end_first = NRA - rows_per_worker*num_workers;
181     printf("first gets most: %zu\n", row_end_first);
182
183     //setup requests
184     MPI_Request requests[num_workers];
185     MM_input *data_first = (MM_input*)malloc(sizeof(MM_input));
186     data_first->rows = row_end_first;
187     data_first->a = (double*)a; //they both start of with no offset!
188     data_first->b = (double*)b_transposed;
189     size_t total_size = sizeof(size_t) + (data_first->rows * NCA)*sizeof(double)+SIZE_OF_B;
190     char* buffer = getbuffer(data_first, total_size); //first one
191
192     // Tag is just nr-cpu -1
193     MPI_Isend(buffer, total_size, MPI_CHAR, 1, 0, MPI_COMM_WORLD, &requests[0]);
194     free(data_first);
195     total_size = sizeof(size_t) + (rows_per_worker * NCA)*sizeof(double) + SIZE_OF_B; //size

```

```

194 is the same for all other - just compute once!
195 size_t i;
196 for (i = 0; i < (num_workers-1); ++i)
197 {
198     MM_input *data = (MM_input*)malloc(sizeof(MM_input));
199     data->rows = rows_per_worker;
200     data->a = (double*)(a + (row_end_first + rows_per_worker*i));
201     data->b = (double*)(b_transposed); // send everything - all needed
202     buffer = getbuffer(data, total_size);
203     printf("nr_worker - %zu\n", i);
204     MPI_Isend(buffer, total_size, MPI_CHAR, i+2, i+1, MPI_COMM_WORLD, &requests[i+1]);
205     free(data);
206 }
207 double* my_a = (double*)(a + (row_end_first + rows_per_worker*i));
208 //I multiply the rest
209 size_t offset = 0;
210 for (size_t row = (NRA-rows_per_worker); row < NRA; row++)
211 {
212     for (size_t col = 0; col < NCB; col++)
213     {
214         res[row * NCB + col] = multsum(my_a+offset, (((double*)b_transposed)+col*NCA), NCA);
215     }
216     offset += NCA;
217 }
218 printf("My c: \n");
219 //wait for rest
220 MPI_Status stats[num_workers];
221 if(MPI_Waitall(num_workers, requests, stats) == MPI_ERR_IN_STATUS){
222     printf("Communication failed!!! - abort\n");
223 }
224 printf(">>>Everything sent and recieved\n");
225
226 // reviece rest
227 size_t buf_size = sizeof(double)*row_end_first*NCB;
228 double* revbuf;
229 offset = 0;
230 for (size_t worker = 0; worker < num_workers; worker++)
231 {
232     revbuf = (double*)malloc(buf_size); //first gets largest buffer
233     MPI_Recv(revbuf, buf_size/sizeof(double), MPI_DOUBLE, worker+1, worker, MPI_COMM_WORLD,
234     &stats[worker]);
235     memcpy(&res[offset/sizeof(double)], revbuf, buf_size);
236     free(revbuf);
237     offset += buf_size;
238     buf_size = sizeof(double)*rows_per_worker*NCB;
239 }
240 double time = MPI_Wtime()-start_time;
241 //free all pointers!
242 free(a);
243 free(b);
244 free(b_transposed);
245 free(c);
246 return time;
247 }
248
249
250 double work(int rank, size_t num_workers){
251     size_t rows_per_worker = NRA / (num_workers+1);
252     char* buffer;
253     MPI_Status status;
254     if (rank == 1) // first always get's most work
255     {
256         rows_per_worker = NRA - rows_per_worker*num_workers;
257     }
258     size_t size_of_meta = sizeof(size_t);
259     size_t size_of_a = sizeof(double)*rows_per_worker*NCA;
260     size_t buffersize = size_of_meta+size_of_a + SIZE_OF_B;
261     buffer = (char*)malloc(buffersize);
262
263     MPI_Recv(buffer, buffersize, MPI_CHAR, 0, rank-1, MPI_COMM_WORLD, &status);

```



```

264     double start = MPI_Wtime();
265     int count;
266     MPI_Get_count(&status, MPI_CHAR, &count);
267     printf("I'm rank %d and I got %d bytes (%ld doubles) of data from %d with tag %d.\n", rank
, count, (count-sizeof(size_t))/sizeof(double), status.MPI_SOURCE, status.MPI_TAG);
268
269     MM_input *mm = (MM_input*)malloc(sizeof(MM_input));
270     mm->a = (double*)&buffer[size_of_meta];
271     mm->b = (double*)&buffer[size_of_meta+size_of_a];
272
273     double *res =(double*)malloc(sizeof(double)*rows_per_worker*NCB);
274
275     size_t offset = 0;
276     for (size_t row = 0; row < rows_per_worker; row++)
277     {
278         for (size_t col = 0; col < NCB; col++)
279         {
280             res[row * NCB + col] = multsum(mm->a+offset, (((double*)mm->b)+col*NCA), NCA);
281         }
282         offset += NCA;
283     }
284     MPI_Send(res, rows_per_worker*NCB, MPI_DOUBLE, 0,rank-1, MPI_COMM_WORLD);
285     printf("[%d] sent res home\n",rank);
286     free(res);
287     return MPI_Wtime() - start;
288 }
289
290 int main(int argc, char *argv[]) {
291     int tid, nthreads;
292     /* for simplicity, set NRA=NCA=NCB=N */
293     // Initialize MPI, find out MPI communicator size and process rank
294     int myRank, numProcs;
295     MPI_Status status;
296     MPI_Init(&argc, &argv);
297     MPI_Comm_size(MPI_COMM_WORLD, &numProcs);
298     MPI_Comm_rank(MPI_COMM_WORLD, &myRank);
299     int num_Workers = numProcs-1;
300     if (argc > 1 && strcmp(argv[1], "parallel") == 0) {
301         // Variables for the process rank and number of processes
302         if (myRank == 0) {
303             printf("Run parallel!\n");
304             double *truth = malloc(sizeof(double) * NRA * NCB);
305             double time = productSequential(truth);
306             printf("Computed reference results in %.6f s\n", time);
307             printf("Hello from master! - I have %d workers!\n", num_Workers);
308             // send out work
309             double *res = malloc(sizeof(double)*NRA*NCB);
310             time = splitwork(res, num_Workers);
311             if (checkResult(res, truth, NCB, NRA)) {
312                 printf("Matrices do not match!!!\n");
313                 return 1;
314             }
315             printf("Matrices match (parallel [eps %.10f])! - took: %.6f s\n", EPS, time);
316             free(truth);
317             free(res);
318         } else {
319             double time = work(myRank, num_Workers);
320             printf("Worker bee %d took %.6f s (after recv) for my work\n", myRank, time);
321         }
322     } else // run sequential
323     {
324         printf("Run sequential!\n");
325         double *res = malloc(sizeof(double) * NRA * NCB);
326         double time = productSequential(res);
327         if (checkResult(res, res, NCB, NRA)) {
328             printf("Matrices do not match!!!\n");
329             return 1;
330         }
331         printf("Matrices match (sequential-trivial)! - took: %.6f s\n", time);
332         free(res);
333     }
334 }
335

```

```

336     MPI_Finalize();
337     return 0;
338 }

```

Appendix - Poisson solver

The parallel Poisson solver used the following code:

```

1  /*
2  * MPI_Poisson.c
3  * 2D Poisson equation solver (parallel version)
4  */
5
6  #include <stdio.h>
7  #include <stdlib.h>
8  #include <math.h>
9  #include <time.h>
10 #include <mpi.h>
11 #include <assert.h>
12
13 #define DEBUG 0
14
15 #define max(a,b) ((a)>(b)?a:b)
16
17
18 // defines for Exercises!
19
20 #define SOR 1
21 #define MONITOR_ERROR 1
22 #define FAST_DO_STEP_LOOP
23 // #define MONITOR_ALLREDUCE 1
24 // #define ALLREDUCE_COUNT 100
25 // #define MONITOR_EXCHANGE_BORDERS
26 #define SKIP_EXCHANGE
27
28 #define DEFINES_ON (SOR || MONITOR_ERROR || 0)
29 //defines end
30
31 enum
32 {
33     X_DIR, Y_DIR
34 };
35
36 // only needed for certain configs!
37 #ifndef SOR
38 double sor_omega = 1.9;
39 #endif
40 #ifdef MONITOR_ERROR
41 double *errors=NULL;
42 #endif
43 #ifdef MONITOR_ALLREDUCE
44 double all_reduce_time = 0;
45 #endif
46 #ifdef MONITOR_EXCHANGE_BORDERS
47 double exchange_time = 0;
48 #endif
49 #ifdef SKIP_EXCHANGE
50 size_t skip_exchange;
51 #endif
52
53 /* global variables */
54 int gridsize[2];
55 double precision_goal; /* precision_goal of solution */
56 int max_iter; /* maximum number of iterations allowed */
57 int P; //total number of processes
58 int P_grid[2]; // process grid dimensions
59 MPI_Comm grid_comm; //grid communicator
60 MPI_Status status;
61 double hx, hy;
62
63 /* process specific globals*/
64 int proc_rank;

```

```

65 double wtime;
66 int proc_coord[2]; // coords of current process in processgrid
67 int proc_top, proc_right, proc_bottom, proc_left; // ranks of neighboring procs
68 // step 7
69 int offset[2] = {0,0};
70 // step 8
71 MPI_Datatype border_type[2];
72
73 /* benchmark related variables */
74 clock_t ticks; // number of systemticks */
75 int timer_on = 0; // is timer running? */
76
77 /* local grid related variables */
78 double **phi; // grid */
79 int **source; // TRUE if subgrid element is a source */
80 int dim[2]; // grid dimensions */
81
82 void Setup_Grid();
83 double Do_Step(int parity);
84 void Solve();
85 void Write_Grid();
86 void Clean_Up();
87 void Debug(char *mesg, int terminate);
88 void start_timer();
89 void resume_timer();
90 void stop_timer();
91 void print_timer();
92
93 void start_timer()
94 {
95     if (!timer_on){
96         MPI_Barrier(grid_comm);
97         ticks = clock();
98         wtime = MPI_Wtime();
99         timer_on = 1;
100     }
101 }
102
103 void resume_timer()
104 {
105     if (!timer_on){
106         ticks = clock() - ticks;
107         wtime = MPI_Wtime() - wtime;
108         timer_on = 1;
109     }
110 }
111
112 void stop_timer()
113 {
114     if (timer_on){
115         ticks = clock() - ticks;
116         wtime = MPI_Wtime() - wtime;
117         timer_on = 0;
118     }
119 }
120
121 void print_timer()
122 {
123     if (timer_on){
124         stop_timer();
125         printf("(i) Elapsed Wtime %14.6f s (%5.1f%% CPU)\n", proc_rank, wtime, 100.0 * ticks
126             * (1.0 / CLOCKS_PER_SEC) / wtime);
127         resume_timer();
128     }
129     else{
130         printf("(i) Elapsed Wtime %14.6f s (%5.1f%% CPU)\n", proc_rank, wtime, 100.0 * ticks
131             * (1.0 / CLOCKS_PER_SEC) / wtime);
132     }
133 }
134
135 void Debug(char *mesg, int terminate)
136 {
137     if (DEBUG || terminate){

```

```

136     printf("%s\n", mesg);
137 }
138 if (terminate){
139     exit(1);
140 }
141 }
142
143 void Setup_Proc_Grid(int argc, char **argv){
144     int wrap_around[2];
145     int reorder;
146
147     Debug("My_MPI_Init",0);
148
149     // num of processes
150     MPI_Comm_size(MPI_COMM_WORLD, &P);
151
152     //calculate the number of processes per column and per row for the grid
153     if(argc>2){
154         P_grid[X_DIR] = atoi(argv[1]);
155         P_grid[Y_DIR] = atoi(argv[2]);
156         if(P_grid[X_DIR] * P_grid[Y_DIR] != P){
157             Debug("ERROR Proces grid dimensions do not match with P ", 1);
158         }
159         #ifdef SOR
160         if (argc>3)
161         {
162             // get sor from args
163             sor_omega = atof(argv[3]);
164             printf("Set sor_omega over argv to %.4f\n", sor_omega);
165         }
166         #endif
167         #ifdef SKIP_EXCHANGE
168         if (argc > 4)
169         {
170             skip_exchange = atoi(argv[4]);
171             printf("Set skip_exchange over argv to %zu\n", skip_exchange);
172         }
173         else{
174             skip_exchange = 1;
175             printf("Set skip_exchange to default value 1\n");
176         }
177         #endif
178     }
179     else{
180         Debug("ERROR Wrong parameter input",1);
181     }
182
183     // Create process topology (2D grid)
184     wrap_around[X_DIR] = 0;
185     wrap_around[Y_DIR] = 0;
186     reorder = 1; //reorder process ranks
187
188     // create grid_comm
189     int ret = MPI_Cart_create(MPI_COMM_WORLD, 2, P_grid, wrap_around, reorder, &grid_comm);
190     if (ret != MPI_SUCCESS){
191         Debug("ERROR: MPI_Cart_create failed",1);
192     }
193     //get new rank and cartesian coords of this proc
194     MPI_Comm_rank(grid_comm, &proc_rank);
195     MPI_Cart_coords(grid_comm, proc_rank, 2, proc_coord);
196     printf("(%) (x,y)=(%,%)\n", proc_rank, proc_coord[X_DIR], proc_coord[Y_DIR]);
197     //calc neighbours
198     // MPI_Cart_shift(grid_comm, Y_DIR, 1, &proc_bottom, &proc_top);
199     MPI_Cart_shift(grid_comm, Y_DIR, 1, &proc_top, &proc_bottom);
200     MPI_Cart_shift(grid_comm, X_DIR, 1, &proc_left, &proc_right);
201     printf("(%) top %, right %, bottom %, left %\n", proc_rank, proc_top,
202     proc_right, proc_bottom, proc_left);
203 }
204
205 void Setup_Grid()
206 {
207     int x, y, s;
208     double source_x, source_y, source_val;

```

```

208 FILE *f;
209
210 Debug("Setup_Subgrid", 0);
211
212 if(proc_rank == 0){
213     f = fopen("input.dat", "r");
214     if (f == NULL){
215         Debug("Error opening input.dat", 1);
216     }
217     fscanf(f, "nx: %i\n", &gridsize[X_DIR]);
218     fscanf(f, "ny: %i\n", &gridsize[Y_DIR]);
219     fscanf(f, "precision goal: %lf\n", &precision_goal);
220     fscanf(f, "max iterations: %i\n", &max_iter);
221 }
222 MPI_Bcast(&gridsize, 2, MPI_INT, 0, grid_comm);
223 MPI_Bcast(&precision_goal, 1, MPI_DOUBLE, 0, grid_comm);
224 MPI_Bcast(&max_iter, 1, MPI_INT, 0, grid_comm);
225 hx = 1 / (double)gridsize[X_DIR];
226 hy = 1 / (double)gridsize[Y_DIR];
227
228 /* Calculate dimensions of local subgrid */ //! We do that later now!
229 // dim[X_DIR] = gridsize[X_DIR] + 2;
230 // dim[Y_DIR] = gridsize[Y_DIR] + 2;
231
232 //! Step 7
233 int upper_offset[2] = {0,0};
234 // Calculate top left corner coordinates of local grid
235 offset[X_DIR] = gridsize[X_DIR] * proc_coord[X_DIR] / P_grid[X_DIR];
236 offset[Y_DIR] = gridsize[Y_DIR] * proc_coord[Y_DIR] / P_grid[Y_DIR];
237 upper_offset[X_DIR] = gridsize[X_DIR] * (proc_coord[X_DIR] + 1) / P_grid[X_DIR];
238 upper_offset[Y_DIR] = gridsize[Y_DIR] * (proc_coord[Y_DIR] + 1) / P_grid[Y_DIR];
239
240 // dimensions of local grid
241 dim[X_DIR] = upper_offset[X_DIR] - offset[X_DIR];
242 dim[Y_DIR] = upper_offset[Y_DIR] - offset[Y_DIR];
243 // Add space for rows/columns of neighboring grid
244 dim[X_DIR] += 2;
245 dim[Y_DIR] += 2;
246 //! Step 7 end
247
248 /* allocate memory */
249 if ((phi = malloc(dim[X_DIR] * sizeof(*phi))) == NULL){
250     Debug("Setup_Subgrid : malloc(phi) failed", 1);
251 }
252 if ((source = malloc(dim[X_DIR] * sizeof(*source))) == NULL){
253     Debug("Setup_Subgrid : malloc(source) failed", 1);
254 }
255 if ((phi[0] = malloc(dim[Y_DIR] * dim[X_DIR] * sizeof(**phi))) == NULL){
256     Debug("Setup_Subgrid : malloc(*phi) failed", 1);
257 }
258 if ((source[0] = malloc(dim[Y_DIR] * dim[X_DIR] * sizeof(**source))) == NULL){
259     Debug("Setup_Subgrid : malloc(*source) failed", 1);
260 }
261 for (x = 1; x < dim[X_DIR]; x++)
262 {
263     phi[x] = phi[0] + x * dim[Y_DIR];
264     source[x] = source[0] + x * dim[Y_DIR];
265 }
266
267 /* set all values to '0' */
268 for (x = 0; x < dim[X_DIR]; x++){
269     for (y = 0; y < dim[Y_DIR]; y++)
270     {
271         phi[x][y] = 0.0;
272         source[x][y] = 0;
273     }
274 }
275 /* put sources in field */
276 do{
277     if (proc_rank==0)
278     {
279         s = fscanf(f, "source: %lf %lf %lf\n", &source_x, &source_y, &source_val);
280     }

```

```

281     MPI_Bcast(&s, 1, MPI_INT, 0, grid_comm);
282     if (s==3){
283         MPI_Bcast(&source_x, 1, MPI_DOUBLE, 0, grid_comm);
284         MPI_Bcast(&source_y, 1, MPI_DOUBLE, 0, grid_comm);
285         MPI_Bcast(&source_val, 1, MPI_DOUBLE, 0, grid_comm);
286         x = source_x * gridsize[X_DIR];
287         y = source_y * gridsize[Y_DIR];
288         x = x + 1 - offset[X_DIR]; // Step 7 --> local grid transform
289         y = y + 1 - offset[Y_DIR]; // Step 7 --> local grid transform
290         if(x > 0 && x < dim[X_DIR] - 1 && y > 0 && y < dim[Y_DIR]-1){ // check if in local
grid
291             phi[x][y] = source_val;
292             source[x][y] = 1;
293         }
294     }
295 }
296 while (s==3);
297
298 if(proc_rank==0){
299     fclose(f);
300 }
301 }
302
303 void Setup_MPI_Datatypes()
304 {
305     Debug("Setup_MPI_Datatypes",0);
306
307     // vertical data exchange (Y_Dir)
308     MPI_Type_vector(dim[X_DIR] - 2, 1, dim[Y_DIR], MPI_DOUBLE, &border_type[Y_DIR]);
309     // horizontal data exchange (X_Dir)
310     MPI_Type_vector(dim[Y_DIR] - 2, 1, 1, MPI_DOUBLE, &border_type[X_DIR]);
311
312     MPI_Type_commit(&border_type[Y_DIR]);
313     MPI_Type_commit(&border_type[X_DIR]);
314 }
315
316 int Exchange_Borders()
317 {
318     #ifdef MONITOR_EXCHANGE_BORDERS
319     double time_ = MPI_Wtime();
320     #endif
321     Debug("Exchange_Borders",0);
322     // top direction
323     MPI_Sendrecv(&phi[1][1], 1, border_type[Y_DIR], proc_top, 0, &phi[1][dim[Y_DIR] - 1], 1,
border_type[Y_DIR], proc_bottom, 0, grid_comm, &status);
324     // bottom direction
325     MPI_Sendrecv(&phi[1][dim[Y_DIR] - 2], 1, border_type[Y_DIR], proc_bottom, 0, &phi[1][0],
1, border_type[Y_DIR], proc_top, 0, grid_comm, &status);
326     // left direction
327     MPI_Sendrecv(&phi[1][1], 1, border_type[X_DIR], proc_left, 0, &phi[dim[X_DIR]-1][1], 1,
border_type[X_DIR], proc_right, 0, grid_comm, &status);
328     // right direction
329     MPI_Sendrecv(&phi[dim[X_DIR]-2][1], 1, border_type[X_DIR], proc_right, 0, &phi[0][1], 1,
border_type[X_DIR], proc_left, 0, grid_comm, &status);
330     #ifdef MONITOR_EXCHANGE_BORDERS
331     exchange_time += MPI_Wtime() - time_;
332     #endif
333     return 1;
334 }
335
336 double Do_Step(int parity)
337 {
338     int x, y;
339     double old_phi, c_ij;
340     double max_err = 0.0;
341
342     #ifdef FAST_DO_STEP_LOOP
343     int start_y;
344     for (x = 1; x < dim[X_DIR] - 1; x++){
345         start_y = ((1 + x + offset[X_DIR] + offset[Y_DIR]) % 2 == parity) ? 1 : 2;
346         for (y = start_y; y < dim[Y_DIR] - 1; y += 2){
347             if (source[x][y] != 1){
348                 old_phi = phi[x][y];

```

```

349         #ifndef SOR
350         phi[x][y] = (phi[x + 1][y] + phi[x - 1][y] + phi[x][y + 1] + phi[x][y - 1]) *
0.25;
351         #endif
352         #ifdef SOR
353         c_ij = (phi[x + 1][y] + phi[x - 1][y] + phi[x][y + 1] + phi[x][y - 1] + hx*hy*
source[x][y]) * 0.25 - phi[x][y];
354         phi[x][y] += sor_omega*c_ij;
355         #endif
356         if (max_err < fabs(old_phi - phi[x][y])){
357             max_err = fabs(old_phi - phi[x][y]);
358         }
359     }
360 }
361 }
362 return max_err;
363 #endif
364
365 #ifndef FAST_DO_STEP_LOOP
366 /* calculate interior of grid */
367 for (x = 1; x < dim[X_DIR] - 1; x++){
368     for (y = 1; y < dim[Y_DIR] - 1; y++){
369         if ((x + offset[X_DIR] + y + offset[Y_DIR]) % 2 == parity && source[x][y] != 1){
370             old_phi = phi[x][y];
371             #ifndef SOR
372             phi[x][y] = (phi[x + 1][y] + phi[x - 1][y] + phi[x][y + 1] + phi[x][y - 1]) *
0.25;
373             #endif
374             #ifdef SOR
375             c_ij = (phi[x + 1][y] + phi[x - 1][y] + phi[x][y + 1] + phi[x][y - 1] + hx*hy*
source[x][y]) * 0.25 - phi[x][y];
376             phi[x][y] += sor_omega*c_ij;
377             #endif
378             if (max_err < fabs(old_phi - phi[x][y])){
379                 max_err = fabs(old_phi - phi[x][y]);
380             }
381         }
382     }
383 }
384 return max_err;
385 #endif
386 }
387
388 void Solve()
389 {
390     int count = 0;
391     double delta;
392     double global_delta;
393     double delta1, delta2;
394
395     Debug("Solve", 0);
396
397     /* give global_delta a higher value then precision_goal */
398     global_delta = 2 * precision_goal;
399
400     while (global_delta > precision_goal && count < max_iter)
401     {
402         Debug("Do_Step 0", 0);
403         delta1 = Do_Step(0);
404         #ifndef SKIP_EXCHANGE
405         if (count % skip_exchange == 0 && Exchange_Borders()) // use short circuit evaluation
406             #endif
407         #ifndef SKIP_EXCHANGE
408         Exchange_Borders();
409         #endif
410         Debug("Do_Step 1", 0);
411         delta2 = Do_Step(1);
412         #ifndef SKIP_EXCHANGE
413         if (count % skip_exchange == 0 && Exchange_Borders())
414             #endif
415         #ifndef SKIP_EXCHANGE
416         Exchange_Borders();
417         #endif

```

```

418     delta = max(delta1, delta2);
419     #ifdef MONITOR_ALLREDUCE
420     double time_ = MPI_Wtime();
421     #endif
422     #ifdef ALLREDUCE_COUNT
423     if(count % ALLREDUCE_COUNT == 0){
424         MPI_Allreduce(&delta, &global_delta, 1, MPI_DOUBLE, MPI_MAX, grid_comm);
425     }
426     #endif
427     #ifndef ALLREDUCE_COUNT
428     MPI_Allreduce(&delta, &global_delta, 1, MPI_DOUBLE, MPI_MAX, grid_comm);
429     #endif
430     #ifdef MONITOR_ALLREDUCE
431     all_reduce_time += MPI_Wtime() - time_;
432     #endif
433     #ifdef MONITOR_ERROR
434     if (proc_rank == 0)
435     {
436         errors[count] = global_delta;
437     }
438     #endif
439     count++;
440 }
441
442 printf("(%i) Number of iterations : %i\n", proc_rank, count);
443 #ifdef MONITOR_ALLREDUCE
444 printf("(%i) Allreduce time: %14.6f\n", proc_rank, all_reduce_time);
445 #endif
446 #ifdef MONITOR_EXCHANGE_BORDERS
447 printf("(%i) Exchange time: %14.6f\n", proc_rank, exchange_time);
448 #endif
449 }
450
451 double* get_Global_Grid()
452 {
453     Debug("get_Global_Grid", 0);
454     ///!! DEBUG only
455     for (size_t i = 0; i < dim[X_DIR]; i++)
456     {
457         for (size_t j = 0; j < dim[Y_DIR]; j++)
458         {
459             phi[i][j] = proc_rank;
460         }
461     }
462 }
463
464 // only process 0 needs to store all data!
465 double* global_phi = NULL;
466 if (proc_rank == 0) {
467     global_phi = malloc(gridsize[X_DIR] * gridsz[e[Y_DIR] * sizeof(double));
468     if (global_phi == NULL) {
469         Debug("get_Global_Grid : malloc(global_phi) failed", 1);
470     }
471 }
472
473 // copy own part into buffer - flatten!
474 size_t buf_size = (dim[X_DIR] - 2) * (dim[Y_DIR] - 2) * sizeof(double);
475 double* local_phi = malloc(buf_size);
476 int idx = 0;
477 for (int x = 1; x < dim[X_DIR] - 1; x++) {
478     for (int y = 1; y < dim[Y_DIR] - 1; y++) {
479         local_phi[idx++] = phi[x][y];
480     }
481 }
482 printf("I'm proc %d and i have a buffer of size %zu\n", proc_rank, buf_size);
483
484
485 // only proc 0 needs sendcounts and displacements for the gather operation
486 int* sendcounts = NULL;
487 int* displs = NULL;
488 if (proc_rank == 0) {
489     sendcounts = malloc(P * sizeof(int));
490     displs = malloc(P * sizeof(int));

```



```

491 // size and offset of different subgrids
492 //! Note that this only works if every process has the same subgrid
493 if (gridsize[X_DIR] % P_grid[X_DIR] != 0 || gridsize[Y_DIR] % P_grid[Y_DIR] != 0)
494 {
495     Debug("!!!A grid dimension is not a multiple of the P_grid in this direction!", 1)
496 ;
497 }
498
499 int subgrid_width = gridsize[X_DIR] / P_grid[X_DIR];
500 int subgrid_height = gridsize[Y_DIR] / P_grid[Y_DIR];
501 for (int px = 0; px < P_grid[X_DIR]; px++) {
502     for (int py = 0; py < P_grid[Y_DIR]; py++) {
503         int rank = px * P_grid[Y_DIR] + py;
504         sendcounts[rank] = subgrid_width * subgrid_height;
505         displs[rank] = (px * subgrid_width * gridsize[Y_DIR]) + (py * subgrid_height);
506     }
507 }
508
509 Debug("get_Global_Grid : MPI_Gatherv", 0);
510 //! TODO this Gatherv does something wrong - all local grids are alright!!!
511 MPI_Gatherv(local_phi, (dim[X_DIR] - 2) * (dim[Y_DIR] - 2), MPI_DOUBLE, global_phi,
512 sendcounts, displs, MPI_DOUBLE, 0, MPI_COMM_WORLD);
513
514 free(local_phi);
515 if (proc_rank == 0) {
516     free(sendcounts);
517     free(displs);
518 }
519
520 return global_phi;
521 }
522
523 void Write_Grid_global(){
524     int x, y;
525     FILE *f;
526     char filename[40]; //seems danagerous to use a static buffer but let's go with the steps
527     sprintf(filename, "output_MPI_global%i.dat", proc_rank);
528     if ((f = fopen(filename, "w")) == NULL){
529         Debug("Write_Grid : fopen failed", 1);
530     }
531
532     Debug("Write_Grid", 0);
533
534     for (x = 1; x < dim[X_DIR]-1; x++){
535         for (y = 1; y < dim[Y_DIR]-1; y++){
536             int x_glob = x + offset[X_DIR];
537             int y_glob = y + offset[Y_DIR];
538             fprintf(f, "%i %i %f\n", x_glob, y_glob, phi[x][y]);
539         }
540     }
541     fclose(f);
542 }
543
544 void Write_Grid()
545 {
546     double* global_phi = get_Global_Grid();
547     if(proc_rank != 0){
548         assert (global_phi == NULL);
549         return;
550     }
551     int x, y;
552     FILE *f;
553     char filename[40]; //seems danagerous to use a static buffer but let's go with the steps
554     sprintf(filename, "output_MPI%i.dat", proc_rank);
555     if ((f = fopen(filename, "w")) == NULL){
556         Debug("Write_Grid : fopen failed", 1);
557     }
558
559     Debug("Write_Grid", 0);
560
561     for (x = 0; x < gridsize[X_DIR]; x++){
562         for (y = 0; y < gridsize[Y_DIR]; y++){

```

```

562         fprintf(f, "%i %i %f\n", x+1, y+1, global_phi[x*gridsize[Y_DIR] + y]);
563     }
564 }
565 fclose(f);
566 free(global_phi);
567 }
568
569 void Clean_Up()
570 {
571     Debug("Clean_Up", 0);
572
573     free(phi[0]);
574     free(phi);
575     free(source[0]);
576     free(source);
577     #ifdef MONITOR_ERROR
578     free(errors);
579     #endif
580 }
581 void setup_error_monitor(){
582     if (proc_rank != 0)
583     {
584         return;
585     }
586
587     errors = malloc(sizeof(double)*max_iter);
588 }
589 void write_errors(){
590     if(proc_rank != 0){
591         return;
592     }
593     FILE *f;
594     char filename[40]; //seems dangerous to use a static buffer but let's go with the steps
595     sprintf(filename, "errors_MPI.dat");
596     if ((f = fopen(filename, "w")) == NULL){
597         Debug("Write_Errors : fopen failed", 1);
598     }
599
600     Debug("Write_Errors", 0);
601
602     for (size_t i = 0; i < max_iter; ++i)
603     {
604         fprintf(f, "%f\n", errors[i]);
605     }
606     fclose(f);
607 }
608 int main(int argc, char **argv)
609 {
610     MPI_Init(&argc, &argv);
611     Setup_Proc_Grid(argc,argv); // was earlier MPI_Comm_rank(MPI_COMM_WORLD, &proc_rank);
612     start_timer();
613
614     Setup_Grid();
615     Setup_MPI_Datatypes();
616
617     #ifdef SOR
618     if (proc_rank == 0)
619     {
620         printf("SOR using omega: %.5f\n", sor_omega);
621     }
622     #endif
623     #ifdef MONITOR_ERROR
624     setup_error_monitor();
625     #endif
626
627     Solve();
628     #ifdef MONITOR_ERROR
629     write_errors();
630     #endif
631     // Write_Grid();
632     Write_Grid_global();
633     print_timer();
634 }

```

```
635     Clean_Up();  
636     MPI_Finalize();  
637     return 0;  
638 }
```