

DELFT UNIVERSITY OF TECHNOLOGY

INTRODUCTION TO HIGH PERFORMANCE COMPUTING  
WI4049TU

---

# Lab Report

---

*Author:*  
Elias Wachmann (6300421)

November 20, 2024



## General Remarks

This final Lab report includes the answers for the exercises (base grad denoted in paranthesis):

0. Introductory exercise (0.5)
1. Poisson solver (1.75)
2. Finite elements simulation (1.0)
3. Eigenvalue solution by Power Method on GPU (1.75)

The optional **shining points** (e.g., performance analysis, optimization, discussion, and clarifying figures) which yield further points are usually marked by a small blue heading in the text or an additional note is added under a figure or table. For example:

**This is a shining point.**

## 0 Introductory exercise

In the introductory lab session, we are taking a look at some basic features of MPI. We start out very simple with a hello world program on two nodes.

### Hello World

```
1 #include "mpi.h"
2 #include <stdio.h>
3
4 int np, rank;
5
6 int main(int argc, char **argv)
7 {
8     MPI_Init(&argc, &argv);
9     MPI_Comm_size(MPI_COMM_WORLD, &np);
10    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
11
12    printf("Node %d of %d says: Hello world!\n", rank, np);
13
14    MPI_Finalize();
15    return 0;
16 }
```

This program can be compiled with the following command:

```
mpicc -o helloworld1.out helloworld1.c
```

And run with:

```
srunc -n 2 -c 4 --mem-per-cpu=1GB ./helloworld1.out
```

We get the following output:

```
Node 0 of 2 says: Hello world!
Node 1 of 2 says: Hello world!
```

From now on I'll skip the compilation and only mention on how many nodes the program is run and what the output is / interpretation of the output.

### 0.a) Ping Pong

I used the template to check how long `MPI_Send` and `MPI_Recv` take. The code can be found in the appendix for this section.

I've modified the printing a bit to make it easier to gather the information. Then I piped the program output into a textfile for further processing in python. I ran it first on one and then on two nodes as specified in the

assignment sheet. Opposed to the averaging over 5 send / receive pairs, I've done 1000 pairs. Furthermore I reran the whole program 5 times to gather more data. All this data is shown in the following graph:



Figure 1: Ping Pong: Number of bytes sent vs. average time taken from 1000 pairs of send / receive. 5 runs shown for each size as scatter plot. Mean of these 5 runs shown as line. Blue small fit includes all data points up to 131072 bytes, blue large from there. Red small fit includes all data points up to 32768 bytes, red large from there.

As can be seen in the data and the fits, there are outliers especially for the larger data sizes. For our runs we get the following fits and R<sup>2</sup> values:

Run Type	Data Size	Fit Equation	R <sup>2</sup> Value
Single Node	Small ( $\leq 131072$ )	$5.95 \times 10^{-7} \cdot x + 7.97 \times 10^{-4}$	0.92
Single Node	Large ( $\geq 131072$ )	$4.61 \times 10^{-7} \cdot x + 1.23 \times 10^{-2}$	0.89
Two Node	Small ( $\leq 32768$ )	$1.07 \times 10^{-6} \cdot x + 2.60 \times 10^{-3}$	0.97
Two Node	Large ( $\geq 32768$ )	$4.41 \times 10^{-7} \cdot x + 3.42 \times 10^{-3}$	0.97

Table 1: Fit Equations and R<sup>2</sup> Values for Single Node and Two Node Runs

**Note:** Each run was performed 5 times (for 1 and 2 nodes) to get a fit on the data and calculate a R<sup>2</sup> value.

**TODO: Further analysis needed?**

### Extra: Ping Pong with MPI\_SendRecv

We do the same analysis for the changed program utilizing MPI\_SendRecv. The code can be found in the appendix for this section.

We get the following graph from the measurements which were performed in the same way as for the previous program:

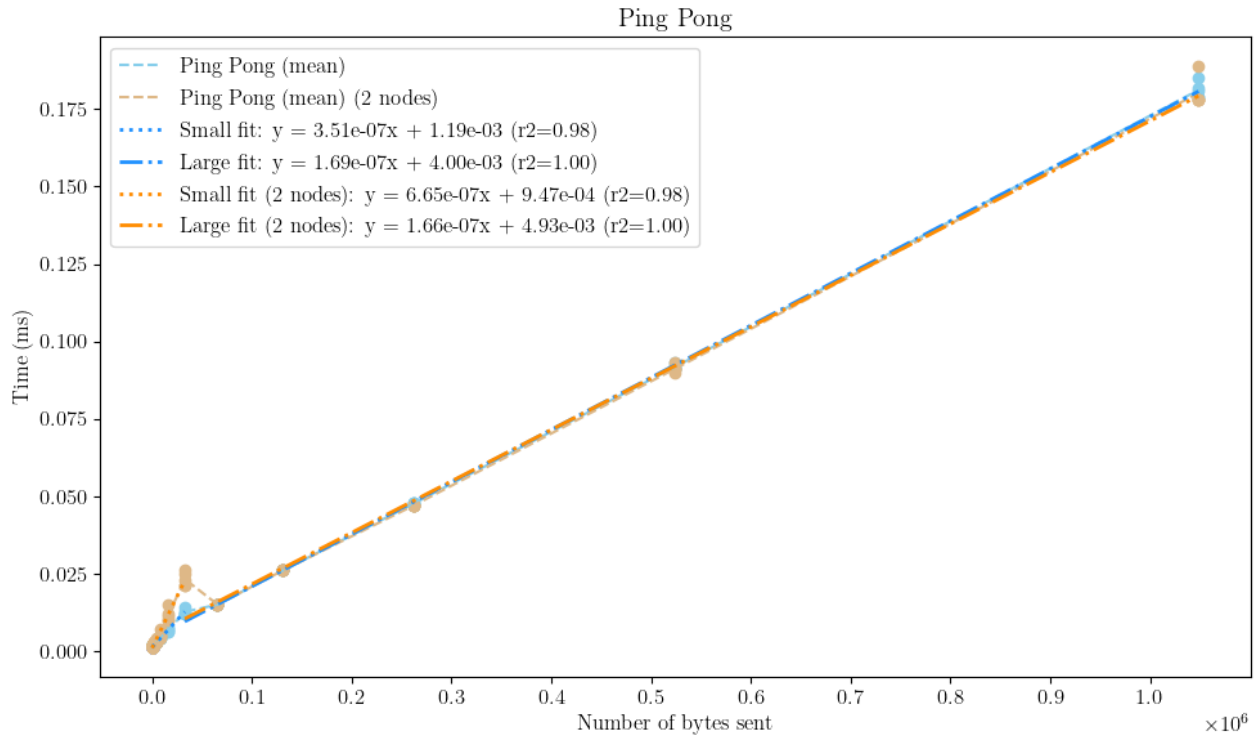


Figure 2: Ping Pong with MPI\_SendRecv: Number of bytes sent vs. average time taken from 1000 pairs of send / receive. 5 runs shown for each size as scatter plot. Mean of these 5 runs shown as line. Blue small fit includes all data points up to 32768 bytes, blue large from there. Red small fit includes all data points up to 32768 bytes, red large from there.

We get the following fits and Rš values for the runs:

Run Type	Data Size	Fit Equation	Rš Value
Single Node	Small ( $\leq 32768$ )	$3.51 \times 10^{-7} \cdot x + 1.19 \times 10^{-3}$	0.98
Single Node	Large ( $\geq 32768$ )	$1.69 \times 10^{-7} \cdot x + 4.00 \times 10^{-3}$	1.00
Two Node	Small ( $\leq 32768$ )	$6.65 \times 10^{-7} \cdot x + 9.47 \times 10^{-4}$	0.98
Two Node	Large ( $\geq 32768$ )	$1.66 \times 10^{-7} \cdot x + 4.93 \times 10^{-3}$	1.00

Table 2: Fit Equations and Rš Values for Single Node and Two Node Runs

**TODO: Further analysis needed?**

## 0.b) MM-product

After an introduction of the matrix-matrix multiplication code in the next section, the measured speedups are discussed in the subsequent section.

### Explanation of the code

For this exercise I've used the template provided in the assignment sheet as a base to develop my parallel implementation for a matrix-matrix multiplication. The code can be found in the appendix for this section.

The program can be run either in sequential (default) or parallel mode (parallel as a command line argument). For the sequential version, the code is practically unchanged and just refactored into a function for timing purposes. The parallel version is more complex and works as explained below:

First, rank 0 computes a sequential reference solution. Then rank 0 distributes the matrices in the following way in `splitwork`:

- Matrix A is split row-wise by dividing the number of rows by the number of nodes.
- The first worker (=rank 1) gets the most rows starting from row 0:  
 $\text{total\_rows} - (\text{nr\_workers} - 1) \cdot \text{floor}(\frac{\text{total\_rows}}{\text{nr\_workers}})$ .
- All other workers and the master (= rank 0) get the same number of rows:  $\text{floor}(\frac{\text{total\_rows}}{\text{nr\_workers}})$ .
- The master copies the corresponding rows of matrix A and the whole transposed matrix B\* into a buffer (for details on MM\_input buffer see below) for each worker and sends them off using MPI\_Isend.
- The workers receive the data using MPI\_Recv and then compute their part of the matrix product and send only the rows of the result matrix back to the master using MPI\_Send.
- In the meanwhile the master computes its part of the matrix product.
- Using MPI\_Waitall the master waits for all data to be sent to the workers and only afterwards calls MPI\_Recv to gather the results from the workers.
- Finally all results are gathered by the master in the result matrix.

Assume we have a 5x5 matrix A and 2 workers (rank 1 and rank 2) and master (rank 0). The partitioning is done row-wise as follows:

#### Partitioning Example

$$A = \begin{pmatrix} a_{11} & a_{12} & a_{13} & a_{14} & a_{15} \\ a_{21} & a_{22} & a_{23} & a_{24} & a_{25} \\ a_{31} & a_{32} & a_{33} & a_{34} & a_{35} \\ a_{41} & a_{42} & a_{43} & a_{44} & a_{45} \\ a_{51} & a_{52} & a_{53} & a_{54} & a_{55} \end{pmatrix} \rightarrow \begin{pmatrix} \text{Worker 1} \\ \text{Worker 1} \\ \text{Worker 1} \\ \text{Master} \\ \text{Master} \end{pmatrix}$$

- **Rank 0 (Master):** Rows 4 and 5 (last two rows)
- **Rank 1 (Worker 1):** Rows 1 to 3 (first three rows) - Worker 1 always gets the most rows

This partitioning can be visually represented as:

$$\text{Master (rank 0): } \begin{pmatrix} a_{41} & a_{42} & a_{43} & a_{44} & a_{45} \\ a_{51} & a_{52} & a_{53} & a_{54} & a_{55} \end{pmatrix}$$

$$\text{Worker 1 (rank 1): } \begin{pmatrix} a_{11} & a_{12} & a_{13} & a_{14} & a_{15} \\ a_{21} & a_{22} & a_{23} & a_{24} & a_{25} \\ a_{31} & a_{32} & a_{33} & a_{34} & a_{35} \end{pmatrix}$$

Each worker computes its part of the matrix product, and the master gathers the results at the end and compiles them into the final matrix.

The MM\_input buffer is used to store the rows of matrix A and the whole matrix B for each worker. It is implemented using a simple struct:

```
1 typedef struct MM_input {
2     size_t rows;
3     double *a;
4     double *b;
5 } MM_input;
```

**\*[Optimization] Note on transposed matrix B:** It is usually beneficial from a cache perspective to index arrays sequentially or in a row-major order. However, in the matrix-matrix multiplication, we access the elements of matrix B in a column-wise order. This leads to cache misses and is not optimal. To mitigate this, we can transpose matrix B and then access it in a row-wise order. This is done in the code by the master before sending the data to the workers.

### Discussion of the speedups

The code was run on Delft's cluster with 1, 2, 4, 8, 16, 24, 32, 48, and 64 nodes. For the experiments the matrix size of  $A$  and  $B$  was set to  $2000 \times 2000$ . This means that the program has to evaluate 2000 multiplications and 1999 additions for each element of the resulting matrix  $C$ . In total this results in  $\approx 2000^3 = 8 \times 10^9$  operations. The command looked similar to the following for the different node counts:

```
srun -n 48 --mem-per-cpu=4GB --time=00:02:00 ./MM.out parallel
```

For this experiment, the execution time was measured and the speedup was calculated. The results are shown in [Table 3](#) and [Figure 3](#).

CPU Count	Execution Time / s	Approx. Speedup
1	47.11	1.0
2	10.26	4.6
4	10.30	4.6
8	5.20	9.1
16	2.97	15.9
24	2.54	18.5
32	2.29	20.6
48	2.98	15.8
64	1.72	27.4

Table 3: Execution Time vs CPU Count



Figure 3: Speedup vs CPU Count  
Black  $\times$  marks the average of the rerun for  $n = 48$ .

**Note:** The speedup is calculated as  $S = \frac{T_1}{T_p}$ , where  $T_1$  is the execution time on 1 node and  $T_p$  is the execution time on  $p$  nodes.

### Discussion:

As one can clearly discern from the data in [Table 3](#) and [Figure 3](#), the speedup increases with the number of nodes (with the exception of  $n = 48$ ). This is expected as the more nodes we have, the more work can be done in

parallel. However, the speedup is not linear. This is due to the overhead of communication between the nodes. The more nodes we have, the more communication is needed, and this overhead increases. This is especially visible in the data for  $n = 48$ . Here the speedup is lower than for  $n = 32$ . For this run the communication didn't went as smooth as for the other runs. This can potentially be attributed to the fact that one (or more) of the nodes or the network was under heavy load during this task.

**[Further investigation]** After observing this slower speed for the  $n = 48$ , I reran the tests multiple times and got a runtime of around 1.9s which was to be expected initially. Therefore, this one run is an odd one out, most likely due to the reasons mentioned above! I've also added the averaged data of the reruns as a datapoint in [Figure 3](#).

Another interesting fact can be seen when comparing the time taken for  $n = 1$  and  $n = 2$ . They don't at all scale with the expected factor of 2. This is could be due to the fact, that the resource management system prefers runs with multiple nodes instead of a single node (= sequential).

Additional notes: The flag `-mem-per-cpu=<#>GB` was set depending on the number of nodes used. For 1-24 nodes 8GB was used, for 32-48 nodes 4GB, and for 64 nodes 3GB. This had to be done to comply with QOS policy on the cluster.

**TODO: Data locality?**

# 1 Poisson solver

In this section of the lab report, we will discuss a parallel implementation of the Poisson solver. The Poisson solver is a numerical method used to solve the Poisson equation, which is a partial differential equation that is useful in many areas of physics.

**Note:** For local testing and development I'll run the code with `mpirun` instead of the `srun` command on the cluster.

## 1.1 Building a parallel Poisson solver

For the first part of the exercise we follow the steps lined out in the assignment sheet. I'll comment on the steps 1 through 10 and related questions below. The finished implementation can be found in the appendix for this section.

1. **Step:** After adding `MPI_Init` and `MPI_Finalize`, we can run the program with multiple processes. We can see that the program runs with 4 processes in [Figure 4](#) via the quadrupled output.

```
etschgi1@Deep-Thought:~/REPOS/HPC/01_lab1/src$ mpirun -np 4 ./mpi.out
Number of iterations : 2355
Number of iterations : 2355
Number of iterations : 2355
Elapsed procestime:      0.133189 s
Number of iterations : 2355
Elapsed procestime:      0.134150 s
Elapsed procestime:      0.134474 s
Elapsed procestime:      0.135356 s
```

Figure 4: MPI\_Poisson after Step 1 - Running with 4 processes

2. **Step:** To see which process is doing what, I included the rank of the process for the print statements as shown in [Figure 5](#).

```
etschgi1@Deep-Thought:~/REPOS/HPC/02_lab1/src$ mpirun -np 4 ./mpi.out
(0) Number of iterations : 2355
(2) Number of iterations : 2355
(0) Elapsed procestime:      0.135963 s
(2) Elapsed procestime:      0.137101 s
(3) Number of iterations : 2355
(3) Elapsed procestime:      0.139614 s
(1) Number of iterations : 2355
(1) Elapsed procestime:      0.142026 s
```

Figure 5: MPI\_Poisson after Step 2 - Running with 4 processes

3. **Step:** Next we define `wtime` as a global double and replace the four utility timing functions with the ones given on Brightspace. A quick verification as shown in [Figure 6](#) shows that the program still runs as expected.

```
etschgi1@Deep-Thought:~/REPOS/HPC/02_lab1/src$ mpirun -np 4 ./mpi.out
(3) Number of iterations : 2355
(1) Number of iterations : 2355
(3) Elapsed Wtime      0.134918 s ( 98.5% CPU)
(1) Elapsed Wtime      0.134459 s ( 98.5% CPU)
(0) Number of iterations : 2355
(2) Number of iterations : 2355
(0) Elapsed Wtime      0.138669 s ( 98.5% CPU)
(2) Elapsed Wtime      0.138910 s ( 98.5% CPU)
```

Figure 6: MPI\_Poisson after Step 3 - Running with 4 processes

4. **Step:** Next we check if two processes indeed give the same output. Both need 2355 iterations to converge and the `diff` command returned no output, which means that the files content is identical.
5. **Step:** Now only the process with rank 0 will read data from files and subsequently broadcast it to the others. Testing this again with 2 processes, we see an empty diff of the output files and the same number of iterations needed to converge.



6. **Step:** We create a cartesian grid of processes using `MPI_Cart_create` and use `MPI_Cart_shift` to find the neighbors of each process. We can see that the neighbors are correctly identified in [Figure 7](#).

```
(0) (x,y)=(0,0)
(0) top 1, right -2, bottom -2, left 2
(1) (x,y)=(0,1)
(1) top -2, right -2, bottom 0, left 3
(2) (x,y)=(1,0)
(2) top 3, right 0, bottom -2, left -2
(3) (x,y)=(1,1)
(3) top -2, right 1, bottom 2, left -2
```

Figure 7: MPI\_Poisson after Step 6 - Running with 4 processes on a 2x2 grid

When there is no neighbor in a certain direction, -2 (or `MPI_PROC_NULL`) is returned.

7. **Step:** We overhaul the setup to get a proper local grid for each process. Furthermore, we only save the relevant source fields in the local grid for each process.

**With for instance 3 processes you should see that 1 or 2 processes do not do any iteration. Do you understand why?**

If we have a look at the input file we see that there are only 3 source fields in the grid. This means that the process that does not have a source field in its local grid will not do any iterations (or only 1). Therefore, if we have 3 processes and the distribution of source fields as given in the input file only 1 process will do iterations if processes are ordered in x-direction and 2 if ordered in y-direction. From this we can conclude that indeed all processes have different local grids and perform different calculations.

```
etschgi@Deep-Thought:~/REPOS/HPC/02_lab1/src$ mpirun -np 3 ./mpi.out 3 1
(0) (x,y)=(0,0)
(0) top -2, right -2, bottom -2, left 1
(1) (x,y)=(1,0)
(1) top -2, right 0, bottom -2, left 2
(2) (x,y)=(2,0)
(2) top -2, right 1, bottom -2, left -2
(0) Number of iterations : 1
(2) Number of iterations : 1
(2) Elapsed Wtime 0.000668 s ( 95.3% CPU)
(0) Elapsed Wtime 0.000917 s ( 95.9% CPU)
(1) Number of iterations : 695
(1) Elapsed Wtime 0.014772 s ( 95.2% CPU)
```

```
etschgi@Deep-Thought:~/REPOS/HPC/02_lab1/src$ mpirun -np 3 ./mpi.out 1 3
(1) (x,y)=(0,1)
(1) top 2, right -2, bottom 0, left -2
(1) Number of iterations : 1
(2) (x,y)=(0,2)
(2) top -2, right -2, bottom 1, left -2
(0) (x,y)=(0,0)
(0) top 1, right -2, bottom -2, left -2
(1) Elapsed Wtime 0.000616 s ( 95.4% CPU)
(0) Number of iterations : 601
(2) Number of iterations : 723
(0) Elapsed Wtime 0.017636 s ( 95.3% CPU)
(2) Elapsed Wtime 0.017801 s ( 95.3% CPU)
```

Figure 8: MPI\_Poisson after Step 7 - Running with 3 processes on a 3x1 (left) vs. 1x3 (right) grid  
For the 3x1 grid, only rank 1 does iterations (> 1), for the 1x3 grid, ranks 0 and 2 do iterations (> 1).

8. **Step:** After defining and committing two special datatypes for vertical and horizontal communication, we setup the communication logic to exchange the boundary values between the processes. We call our `Exchange_Borders` function after each iteration (for both red / black grid points). Now we face the problem in which some processes may stop instantly (no source in their local grid). They will not supply any data to their neighbors, which will cause the program to hang. We shall fix this in the next step.
9. **Step:** Finally we need to implement the logic to check for convergence (in a global sense). We do this by using a `MPI_Allreduce` call with the `MPI_MAX` operation. This way we aggregate all deltas and choose the biggest one for the global delta which we use in the while-loop-condition to check for convergence. We can see that the program now runs as expected in [Figure 9](#).

```

(0) (x,y)=(0,0)
(0) top -1, right 2, bottom 1, left -1
(1) (x,y)=(0,1)
(1) top 0, right 3, bottom -1, left -1
(2) (x,y)=(1,0)
(2) top -1, right -1, bottom 3, left 0
(3) (x,y)=(1,1)
(3) top 2, right -1, bottom -1, left 1
(0) Number of iterations : 2355
(1) Number of iterations : 2355
(2) Number of iterations : 2355
(3) Number of iterations : 2355
(1) Elapsed Wtime      0.287549 s ( 99.9% CPU)
(2) Elapsed Wtime      0.287537 s (100.0% CPU)
(3) Elapsed Wtime      0.287537 s (100.0% CPU)
(0) Elapsed Wtime      0.295957 s ( 99.9% CPU)

```

Figure 9: MPI\_Poisson after Step 9 - Running with 4 processes on a 2x2 grid

Note that this run in Figure 9 was done with another pc and another MPI implementation. Therefore, we see  $-1$  for cells without a neighbor! However, other than that cosmetic difference it has no impact on the programm.

10. **Step:** Now we only have to fix two remaining things. First we have to make sure that each process uses the right global coordinates for the output file in the end. Therefore, we change the function a bit to include the specific x-/y-offset for each processor. The second thing is the potential problem, that different processors might start with different (red/black) parities. In order to accomplish a global parity we simply have to change the calculation in the if in `Do_Step` from

```
1 if ((x + y) % 2 == parity && source[x][y] != 1)
```

to

```
1 if ((x + offset[X_DIR] + y + offset[Y_DIR]) % 2 == parity && source[x][y] != 1)
```

this guarantees that during a given iteration all processors are using the same parity.

This just leaves one question open: Are the results acutally the same?

Checking the output files of the MPI-implementation with the sequential reference indeed shows identical numerical values for the calculated points. Furthermore, the needed iterationcount is also identical which isn't a big surprise, given that the two programmes perform the exact same calculation steps.

## 1.2 Exercises, modifications, and performance aspects

For this subsection we'll define the following shorthand notation:

$n$ :	the number of iterations
$g$ :	gridsize
$t$ :	time needed in seconds
$pt$ :	processor topology in form $pxy$ , where:
$p$ :	number of processors used
$x$ :	number of processors in x-direction
$y$ :	number of processors in y-direction

Table 4: Notation for this section

$pt = 414$  means 4 processors in a  $1 \times 4$  topology.

### 1.2.1 Over-relaxation (SOR)

We start of by rewriting the `Do_Step` routine to facilitate SOR updates. Furthermore, we need  $h^2$ , the grid spacing (which is 1 in our case) and the relaxation parameter  $\omega$  to calculate the updated values. A quick test shows a speedup of roughly a factor of 10. More systematic tests will be done in the next section.

### 1.2.2 Optimal $\omega$ for 4 proc. on a 4x1 grid

With the power of a little python scripting we can easily test different values for  $\omega$  and plot the results as seen in Figure 10.

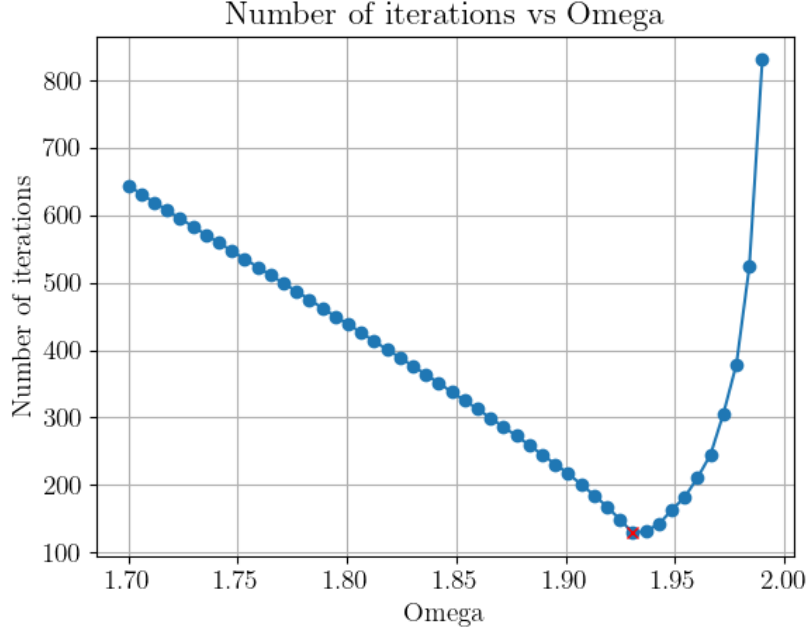


Figure 10: Optimal  $\omega$  for 4 processors on a 4x1 grid

We find that the optimal  $\omega$  is at about 1.93 for this setup with only 129 iterations. This constitutes a speedup of about 1825% compared to the sequential implementation.

### 1.2.3 Scaling behavior with larger grids

This investigation is carried out twice: Once with a  $4 \times 1$  topology (as in the previous section) and once with a  $2 \times 2$  topology. We use grid sizes of  $10 \times 10$ ,  $25 \times 25$ ,  $50 \times 50$ ,  $100 \times 100$ ,  $200 \times 200$ ,  $400 \times 400$ ,  $800 \times 800$  and  $1600 \times 1600$  and set  $\omega = 1.95$  for all runs. The results are shown in Figure 11.

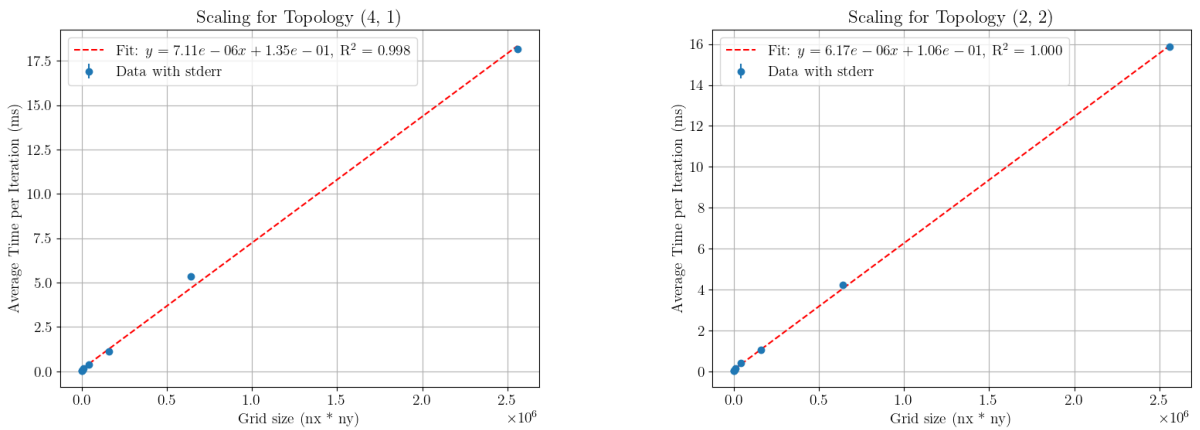


Figure 11: Scaling behavior of the Poisson solver with different grid sizes and processor topologies

As seen by the high  $R^2$  values in the plots, the scaling behavior is very close to linear. We obtain the following scaling factors for the different grid sizes and topologies from the linear fits:

Topology	$\alpha$	$\beta$
$4 \times 1$	$1.35 + 10^{-1}$	$7.11 + 10^{-6}$
$2 \times 2$	$1.06 + 10^{-1}$	$6.17 + 10^{-6}$

Table 5: Scaling factors for different processor topologies for the Poisson solver  
Using:  $t(n) = \alpha + \beta \cdot n$  as a model

#### What can you conclude from the scaling behavior?

We see that the scaling behavior is very close to linear for both topologies. This means that the parallel implementation scales as expected with the number of grid points.

If we compare the scaling factors ( $\beta$ ) for the two topologies we see that the  $2 \times 2$  topology scales slightly better than the  $4 \times 1$  topology. This is not surprising, as the  $2 \times 2$  topology has a more balanced communication workload balance. In the  $2 \times 2$  topology every processor has two neighbors, while in the  $4 \times 1$  topology the processors at the ends only have one neighbor. This is a general trend: A topology which divides the grid into square / square-like parts will scale better than a topology which divides the grid into long and thin parts.

In essence: We want to keep the communication between processors as balanced as possible to achieve the best scaling behavior.

#### 1.2.4 Scaling behavior [Theory - no measurements]

If I could choose between a  $16 \times 1$ ,  $8 \times 2$ ,  $4 \times 4$ ,  $2 \times 8$ ,  $1 \times 16$  topology, I would choose the  $4 \times 4$  topology. This is because the  $4 \times 4$  topology has the most balanced communication workload balance, as detailed in the [Shining](#) in [subsubsection 1.2.3](#).

#### 1.2.5 Iterations needed for convergence scaling

We investigate the number of iterations needed for convergence using the  $4 \times 1$  topology square grids with sidelength: 10, 25, 50, 100, 200, 400, 800, 1600. The results for different  $\omega$  are shown in [Figure 12](#).

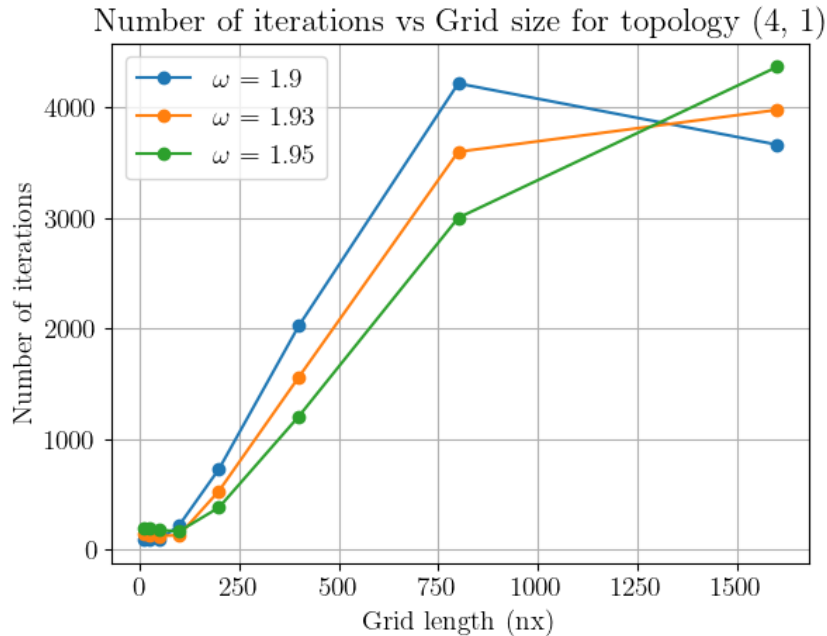


Figure 12: Iterations needed for convergence with different grid side lengths

We can clearly see that the number of iterations till convergence increases with the problem size. At first, I expected linear growth proportional to the number of gridpoints. However, it turns out that the number of iterations actually grow slower and in a square root like fashion. This can be seen by the linear behavior in the plot of grid-side length against iterations.

**Why is the number of iterations needed for convergence  $\propto \sqrt{g}$ ?**

Our poisson problem is a discretized system in 2D space. The condition number of the matrix we have to solve

is proportional to the number of gridpoints in our system. SOR uses the spectral properties of the matrix to solve in a way such that the dominant error mode takes time proportional to the diameter of the domain to converge. This means it is proportional to  $\sqrt{g} = \sqrt{n_x \cdot n_y}$ .

### Why does omega with the best performance change with the grid size?

As can be seen in Figure 12  $\omega = 1.9$  beats the other two values for very small and the largest gridsize. For different gridsizes we get differently sized matrices we have to solve. SOR overrelaxes high-frequency errors and underrelaxes low-frequency errors (the later for stability). The optimal  $\omega$  is therefore dependent on the gridsize and the error modes present in the system. In our current example, it might be that  $\omega = 1.9$  is a good compromise for the grid sizes we are looking at.

#### 1.2.6 Error as a function of the iteration number

With the same  $4 \times 1$  topology and grid sizes of  $800 \times 800$  the error for 15000 iterations is tracked using  $\omega = 1.93$ . The results are shown in Figure 13.




Figure 13: Error as a function of the iteration number

At first the error decreases rapidly in the first few iterations to about  $10^{-3}$  (logarithmic scale!). After that the error decreases more slowly until it is below floating point precision. **Note:** All calculations are done using double precision floating point numbers and only the error recording was done using single precision which leaves the step-like artifacts in the plot.

#### 1.2.7 Optional - Gain performance by reducing MPI\_Allreduce calls

The last subsection showed us that the error reduces monotonically. We might be able to save some time by leaving out some checks and maybe check the global error every 10th or 100th iteration only.

First, we should benchmark if it is at all wise to optimize here, by measuring how long the `MPI_Allreduce` call takes. We can do this by measuring the time needed for the `MPI_Allreduce` call in the `Do_Step` function and summing up to get the total time spent in `MPI_Allreduce` calls.

We again solve with a  $4 \times 1$  topology,  $\omega = 1.93$  and a  $800 \times 800$  grid: It takes roughly 20 seconds of which the processors spend around 1 - 2 seconds in the `MPI_Allreduce` call. This is a significant amount of time ( $(7.0 \pm 0.4)\%$ ). This means we would save some time by reducing the number of `MPI_Allreduce` calls and calculating 9 (0.25% of total) more iterations wouldn't hurt us too bad because it takes 3601 to converge!

We run the program three times with `MPI_Allreduce` calls every 1, 10 and 100 iterations and get the speedups in `MPI_Allreduce` calls as shown in Table 6.

Iterations	MPI_Allreduce - speedup (factor)	calculated overall speedup (%)
1	1.00	0
10	$6.0 \pm 2.0$	$5.9 \pm 0.5$
100	$62 \pm 6$	$6.9 \pm 0.4$

Table 6: Speedup in MPI\_Allreduce calls for different iteration counts and calculated overall speedup (%)

As can be clearly seen from the table we can gain around 6 % using MPI\_Allreduce calls every 10 iterations and around 7 % using MPI\_Allreduce calls every 100 iterations. This is a significant speedup for a very small change in the code.

**Note:** The speedup is calculated to account for fluctuations in the runtime of the program, due to other processes running on the same machine / cluster.

### 1.2.8 Reduce border communication

Another way to reduce communication overhead is to reduce the number of border exchanges.

## 2 Finite elements simulation

### 3 Eigenvalue solution by Power Method on GPU



## Appendix - Introductory exercise

The following code was used for the ping pong task:

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <mpi.h>
4
5 // Maximum array size 2^20= 1048576 elements
6 #define MAX_EXPONENT 20
7 #define MAX_ARRAY_SIZE (1<<MAX_EXPONENT)
8 #define SAMPLE_COUNT 1000
9
10 int main(int argc, char **argv)
11 {
12     // Variables for the process rank and number of processes
13     int myRank, numProcs, i;
14     MPI_Status status;
15
16     // Initialize MPI, find out MPI communicator size and process rank
17     MPI_Init(&argc, &argv);
18     MPI_Comm_size(MPI_COMM_WORLD, &numProcs);
19     MPI_Comm_rank(MPI_COMM_WORLD, &myRank);
20
21
22     int *myArray = (int *)malloc(sizeof(int)*MAX_ARRAY_SIZE);
23     if (myArray == NULL)
24     {
25         printf("Not enough memory\n");
26         exit(1);
27     }
28     // Initialize myArray
29     for (i=0; i<MAX_ARRAY_SIZE; i++)
30         myArray[i]=1;
31
32     int number_of_elements_to_send;
33     int number_of_elements_received;
34
35     // PART C
36     if (numProcs < 2)
37     {
38         printf("Error: Run the program with at least 2 MPI tasks!\n");
39         MPI_Abort(MPI_COMM_WORLD, 1);
40     }
41     double startTime, endTime;
42
43     // TODO: Use a loop to vary the message size
44     for (size_t j = 0; j <= MAX_EXPONENT; j++)
45     {
46         number_of_elements_to_send = 1<<j;
47         if (myRank == 0)
48         {
49             myArray[0]=myArray[1]+1; // activate in cache (avoids possible delay when sending
the 1st element)
50             startTime = MPI_Wtime();
51             for (i=0; i<SAMPLE_COUNT; i++)
52             {
53                 MPI_Send(myArray, number_of_elements_to_send, MPI_INT, 1, 0,
54                     MPI_COMM_WORLD);
55                 MPI_Probe(MPI_ANY_SOURCE, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
56                 MPI_Get_count(&status, MPI_INT, &number_of_elements_received);
57
58                 MPI_Recv(myArray, number_of_elements_received, MPI_INT, 1, 0,
59                     MPI_COMM_WORLD, MPI_STATUS_IGNORE);
60             } // end of for-loop
61
62             endTime = MPI_Wtime();
63             printf("Rank %2.1i: Received %i elements: Ping Pong took %f seconds\n", myRank,
number_of_elements_received, (endTime - startTime)/(2*SAMPLE_COUNT));
64         }
65         else if (myRank == 1)
66         {
67             // Probe message in order to obtain the amount of data
68             MPI_Probe(MPI_ANY_SOURCE, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
```

```

69     MPI_Get_count(&status, MPI_INT, &number_of_elements_received);
70
71     for (i=0; i<SAMPLE_COUNT; i++)
72     {
73         MPI_Recv(myArray, number_of_elements_received, MPI_INT, 0, 0,
74                 MPI_COMM_WORLD, MPI_STATUS_IGNORE);
75         MPI_Send(myArray, number_of_elements_to_send, MPI_INT, 0, 0,
76                 MPI_COMM_WORLD);
77     } // end of for-loop
78 }
79 }
80
81 // Finalize MPI
82 MPI_Finalize();
83
84 return 0;
85 }

```

For the bonus task, the following code was used:

```

1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <mpi.h>
4
5  // Maximum array size 2^20= 1048576 elements
6  #define MAX_EXPONENT 20
7  #define MAX_ARRAY_SIZE (1<<MAX_EXPONENT)
8  #define SAMPLE_COUNT 1000
9
10 int main(int argc, char **argv)
11 {
12     // Variables for the process rank and number of processes
13     int myRank, numProcs, i;
14     MPI_Status status;
15
16     // Initialize MPI, find out MPI communicator size and process rank
17     MPI_Init(&argc, &argv);
18     MPI_Comm_size(MPI_COMM_WORLD, &numProcs);
19     MPI_Comm_rank(MPI_COMM_WORLD, &myRank);
20
21
22     int *myArray = (int *)malloc(sizeof(int)*MAX_ARRAY_SIZE);
23     if (myArray == NULL)
24     {
25         printf("Not enough memory\n");
26         exit(1);
27     }
28     // Initialize myArray
29     for (i=0; i<MAX_ARRAY_SIZE; i++)
30         myArray[i]=1;
31
32     int number_of_elements_to_send;
33     int number_of_elements_received;
34
35     // PART C
36     if (numProcs < 2)
37     {
38         printf("Error: Run the program with at least 2 MPI tasks!\n");
39         MPI_Abort(MPI_COMM_WORLD, 1);
40     }
41     double startTime, endTime;
42
43     // TODO: Use a loop to vary the message size
44     for (size_t j = 0; j <= MAX_EXPONENT; j++)
45     {
46         number_of_elements_to_send = 1<<j;
47         if (myRank == 0)
48         {
49             myArray[0]=myArray[1]+1; // activate in cache (avoids possible delay when sending
the 1st element)
50             startTime = MPI_Wtime();
51             for (i=0; i<SAMPLE_COUNT; i++)
52             {
53                 MPI_Sendrecv(myArray, number_of_elements_to_send, MPI_INT, 1,0,myArray,

```

```

    number_of_elements_to_send, MPI_INT, 1, 0, MPI_COMM_WORLD, &status);
54     }
55
56     endTime = MPI_Wtime();
57     printf("Rank %2.i: Received %i elements: Ping Pong took %f seconds\n", myRank,
    number_of_elements_to_send, (endTime - startTime)/(2*SAMPLE_COUNT));
58     }
59     else if (myRank == 1)
60     {
61         for (i=0; i<SAMPLE_COUNT; i++)
62         {
63             MPI_Sendrecv(myArray, number_of_elements_to_send, MPI_INT, 0,0,myArray,
    number_of_elements_to_send, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
64         }
65     }
66 }
67
68 // Finalize MPI
69 MPI_Finalize();
70
71 return 0;
72 }

```

The matrix multiplication used the following code:

```

1  /*****
2  * FILE: mm.c
3  * DESCRIPTION:
4  *   This program calculates the product of matrix a[nra][nca] and b[nca][ncb],
5  *   the result is stored in matrix c[nra][ncb].
6  *   The max dimension of the matrix is constraint with static array
7  *   declaration, for a larger matrix you may consider dynamic allocation of the
8  *   arrays, but it makes a parallel code much more complicated (think of
9  *   communication), so this is only optional.
10 *
11 *****/
12
13 #include <math.h>
14 #include <mpi.h>
15 #include <stdbool.h>
16 #include <stdio.h>
17 #include <stdlib.h>
18 #include <string.h>
19
20 #define NRA 2000 /* number of rows in matrix A */
21 #define NCA 2000 /* number of columns in matrix A */
22 #define NCB 2000 /* number of columns in matrix B */
23 // #define N 1000
24 #define EPS 1e-9
25 #define SIZE_OF_B NCA*NCB*sizeof(double)
26
27 bool eps_equal(double a, double b) { return fabs(a - b) < EPS; }
28
29 void print_flattened_matrix(double *matrix, size_t rows, size_t cols, int rank) {
30     printf("[%d]\n", rank);
31     for (size_t i = 0; i < rows; i++) {
32         for (size_t j = 0; j < cols; j++) {
33             printf("%10.2f ", matrix[i * cols + j]); // Accessing element in the 1D array
34         }
35         printf("\n"); // Newline after each row
36     }
37 }
38
39 int checkResult(double *truth, double *test, size_t Nr_col, size_t Nr_rows) {
40     for (size_t i = 0; i < Nr_rows; ++i) {
41         for (size_t j = 0; j < Nr_col; ++j) {
42             size_t index = i * Nr_col + j;
43             if (!eps_equal(truth[index], test[index])) {
44                 return 1;
45             }
46         }
47     }
48     return 0;
49 }

```

```

50
51 typedef struct {
52     size_t rows;
53     double *a;
54     double *b;
55 } MM_input;
56
57 char* getbuffer(MM_input *in, size_t size_of_buffer){
58     char* buffer = (char*)malloc(size_of_buffer * sizeof(char));
59     if (buffer == 0)
60     {
61         printf("Buffer couldn't be allocated.");
62         return NULL;
63     }
64     size_t offset = 0;
65     memcpy(buffer + offset, &in->rows, sizeof(size_t));
66     offset += sizeof(size_t);
67     size_t matrix_size = in->rows * NCA * sizeof(double);
68     memcpy(buffer + offset, in->a, matrix_size);
69     offset += matrix_size;
70     memcpy(buffer + offset, in->b, NCA*NCB*sizeof(double));
71     return buffer;
72 }
73
74 MM_input* readbuffer(char* buffer, size_t size_of_buffer){
75     MM_input *mm = (MM_input*)malloc(sizeof(MM_input));
76
77     mm->rows = ((size_t*)buffer)[0];
78     size_t offset = sizeof(size_t);
79     size_t matrix_size = mm->rows * NCA;
80     mm->a = (double*)malloc(sizeof(double)*matrix_size);
81     mm->b = (double*)malloc(sizeof(double)*matrix_size);
82     memcpy(mm->a, &(buffer[offset]), matrix_size);
83     offset += matrix_size;
84     memcpy(mm->b, &(buffer[offset]), NCA*NCB*sizeof(double));
85     free(buffer);
86     return mm;
87 }
88
89
90 void setupMatrices(double (*a)[NCA], double (*b)[NCB], double (*c)[NCB]){
91     for (size_t i = 0; i < NRA; i++) {
92         for (size_t j = 0; j < NCA; j++) {
93             a[i][j] = i + j;
94         }
95     }
96
97     for (size_t i = 0; i < NCA; i++) {
98         for (size_t j = 0; j < NCB; j++) {
99             b[i][j] = i * j;
100         }
101     }
102
103     for (size_t i = 0; i < NRA; i++) {
104         for (size_t j = 0; j < NCB; j++) {
105             c[i][j] = 0;
106         }
107     }
108 }
109
110 double multsum(double* a, double* b_transposed, size_t size){
111     double acc = 0;
112     for (size_t i = 0; i < size; i++)
113     {
114         acc += a[i]*b_transposed[i];
115     }
116     return acc;
117 }
118
119 double productSequential(double *res) {
120     // dynamically allocate to not run into stack overflow - usually stacks are
121     // 8192 bytes big -> 1024 doubles but we have 1 Mio. per matrix
122     double(*a)[NCA] = malloc(sizeof(double) * NRA * NCA);

```

```

123 double(*b)[NCB] = malloc(sizeof(double) * NCA * NCB);
124 double(*c)[NCB] = malloc(sizeof(double) * NRA * NCB);
125
126 /** Initialize matrices */
127 setupMatrices(a,b,c);
128
129 /* Parallelize the computation of the following matrix-matrix
130 multiplication. How to partition and distribute the initial matrices, the
131 work, and collecting final results.
132 */
133 // multiply
134 double start = MPI_Wtime();
135 for (size_t i = 0; i < NRA; i++) {
136     for (size_t j = 0; j < NCB; j++) {
137         for (size_t k = 0; k < NCA; k++) {
138             res[i * NCB + j] += a[i][k] * b[k][j];
139         }
140     }
141 }
142 /* perform time measurement. Always check the correctness of the parallel
143 results by printing a few values of c[i][j] and compare with the
144 sequential output.
145 */
146 double time = MPI_Wtime()-start;
147 free(a);
148 free(b);
149 free(c);
150 return time;
151 }
152
153 double splitwork(double* res, size_t num_workers){
154     if (num_workers == 0) // sadly noone will help me :(
155     {
156         printf("Run sequential!\n");
157         return productSequential(res);
158     }
159
160     double(*a)[NCA] = malloc(sizeof(double) * NRA * NCA);
161     double(*b)[NCB] = malloc(sizeof(double) * NCA * NCB);
162     double(*c)[NCB] = malloc(sizeof(double) * NRA * NCB);
163     // Transpose matrix b to make accessing columns easier - in row major way - better cache
164     // performance
165     setupMatrices(a,b,c);
166
167     double start_time = MPI_Wtime();
168     double (*b_transposed)[NCA] = malloc(sizeof(double) * NCA * NCB);
169     for (size_t i = 0; i < NCA; i++) {
170         for (size_t j = 0; j < NCB; j++) {
171             b_transposed[j][i] = b[i][j];
172         }
173     }
174
175     /** Initialize matrices */
176     // given number of workers I'll split
177     size_t rows_per_worker = NRA / (num_workers+1); //takes corresponding columns from other
178     // matrix
179     printf("rows per worker: %zu\n", rows_per_worker);
180     size_t row_end_first = NRA - rows_per_worker*num_workers;
181     printf("first gets most: %zu\n", row_end_first);
182
183     //setup requests
184     MPI_Request requests[num_workers];
185     MM_input *data_first = (MM_input*)malloc(sizeof(MM_input));
186     data_first->rows = row_end_first;
187     data_first->a = (double*)a; //they both start of with no offset!
188     data_first->b = (double*)b_transposed;
189     size_t total_size = sizeof(size_t) + (data_first->rows * NCA)*sizeof(double)+SIZE_OF_B;
190     char* buffer = getbuffer(data_first, total_size); //first one
191
192     // Tag is just nr-cpu -1
193     MPI_Isend(buffer, total_size, MPI_CHAR, 1, 0, MPI_COMM_WORLD, &requests[0]);
194     free(data_first);
195     total_size = sizeof(size_t) + (rows_per_worker * NCA)*sizeof(double) + SIZE_OF_B; //size

```

```

194 is the same for all other - just compute once!
195 size_t i;
196 for (i = 0; i < (num_workers-1); ++i)
197 {
198     MM_input *data = (MM_input*)malloc(sizeof(MM_input));
199     data->rows = rows_per_worker;
200     data->a = (double*)(a + (row_end_first + rows_per_worker*i));
201     data->b = (double*)(b_transposed); // send everything - all needed
202     buffer = getbuffer(data, total_size);
203     printf("nr_worker - %zu\n", i);
204     MPI_Isend(buffer, total_size, MPI_CHAR, i+2, i+1, MPI_COMM_WORLD, &requests[i+1]);
205     free(data);
206 }
207 double* my_a = (double*)(a + (row_end_first + rows_per_worker*i));
208 //I multiply the rest
209 size_t offset = 0;
210 for (size_t row = (NRA-rows_per_worker); row < NRA; row++)
211 {
212     for (size_t col = 0; col < NCB; col++)
213     {
214         res[row * NCB + col] = multsum(my_a+offset, (((double*)b_transposed)+col*NCA), NCA
215 );
216     }
217     offset += NCA;
218 }
219 printf("My c: \n");
220 //wait for rest
221 MPI_Status stats[num_workers];
222 if(MPI_Waitall(num_workers, requests, stats) == MPI_ERR_IN_STATUS){
223     printf("Communication failed!!! - abort\n");
224 }
225 printf(">>>Everything sent and recieved\n");
226 // reviece rest
227 size_t buf_size = sizeof(double)*row_end_first*NCB;
228 double* revbuf;
229 offset = 0;
230 for (size_t worker = 0; worker < num_workers; worker++)
231 {
232     revbuf = (double*)malloc(buf_size); //first gets largest buffer
233     MPI_Recv(revbuf, buf_size/sizeof(double), MPI_DOUBLE, worker+1, worker, MPI_COMM_WORLD
234 ,&stats[worker]);
235     memcpy(&res[offset/sizeof(double)], revbuf, buf_size);
236     free(revbuf);
237     offset += buf_size;
238     buf_size = sizeof(double)*rows_per_worker*NCB;
239 }
240 double time = MPI_Wtime()-start_time;
241 //free all pointers!
242 free(a);
243 free(b);
244 free(b_transposed);
245 free(c);
246 return time;
247 }
248
249
250 double work(int rank, size_t num_workers){
251     size_t rows_per_worker = NRA / (num_workers+1);
252     char* buffer;
253     MPI_Status status;
254     if (rank == 1) // first always get's most work
255     {
256         rows_per_worker = NRA - rows_per_worker*num_workers;
257     }
258     size_t size_of_meta = sizeof(size_t);
259     size_t size_of_a = sizeof(double)*rows_per_worker*NCA;
260     size_t buffersize = size_of_meta+size_of_a + SIZE_OF_B;
261     buffer = (char*)malloc(buffersize);
262
263     MPI_Recv(buffer, buffersize, MPI_CHAR, 0, rank-1, MPI_COMM_WORLD, &status);

```

```

264     double start = MPI_Wtime();
265     int count;
266     MPI_Get_count(&status, MPI_CHAR, &count);
267     printf("I'm rank %d and I got %d bytes (%ld doubles) of data from %d with tag %d.\n", rank
, count, (count-sizeof(size_t))/sizeof(double), status.MPI_SOURCE, status.MPI_TAG);
268
269     MM_input *mm = (MM_input*)malloc(sizeof(MM_input));
270     mm->a = (double*)&buffer[size_of_meta];
271     mm->b = (double*)&buffer[size_of_meta+size_of_a];
272
273     double *res =(double*)malloc(sizeof(double)*rows_per_worker*NCB);
274
275     size_t offset = 0;
276     for (size_t row = 0; row < rows_per_worker; row++)
277     {
278         for (size_t col = 0; col < NCB; col++)
279         {
280             res[row * NCB + col] = multsum(mm->a+offset, (((double*)mm->b)+col*NCA), NCA);
281         }
282         offset += NCA;
283     }
284     MPI_Send(res, rows_per_worker*NCB, MPI_DOUBLE, 0,rank-1, MPI_COMM_WORLD);
285     printf("[%d] sent res home\n",rank);
286     free(res);
287     return MPI_Wtime() - start;
288 }
289
290 int main(int argc, char *argv[]) {
291     int tid, nthreads;
292     /* for simplicity, set NRA=NCA=NCB=N */
293     // Initialize MPI, find out MPI communicator size and process rank
294     int myRank, numProcs;
295     MPI_Status status;
296     MPI_Init(&argc, &argv);
297     MPI_Comm_size(MPI_COMM_WORLD, &numProcs);
298     MPI_Comm_rank(MPI_COMM_WORLD, &myRank);
299     int num_Workers = numProcs-1;
300     if (argc > 1 && strcmp(argv[1], "parallel") == 0) {
301         // Variables for the process rank and number of processes
302         if (myRank == 0) {
303             printf("Run parallel!\n");
304             double *truth = malloc(sizeof(double) * NRA * NCB);
305             double time = productSequential(truth);
306             printf("Computed reference results in %.6f s\n", time);
307             printf("Hello from master! - I have %d workers!\n", num_Workers);
308             // send out work
309             double *res = malloc(sizeof(double)*NRA*NCB);
310             time = splitwork(res, num_Workers);
311             if (checkResult(res, truth, NCB, NRA)) {
312                 printf("Matrices do not match!!!\n");
313                 return 1;
314             }
315             printf("Matrices match (parallel [eps %.10f])! - took: %.6f s\n", EPS, time);
316             free(truth);
317             free(res);
318         } else {
319             double time = work(myRank, num_Workers);
320             printf("Worker bee %d took %.6f s (after recv) for my work\n", myRank, time);
321         }
322     } else // run sequential
323     {
324         printf("Run sequential!\n");
325         double *res = malloc(sizeof(double) * NRA * NCB);
326         double time = productSequential(res);
327         if (checkResult(res, res, NCB, NRA)) {
328             printf("Matrices do not match!!!\n");
329             return 1;
330         }
331         printf("Matrices match (sequential-trivial)! - took: %.6f s\n", time);
332         free(res);
333     }
334 }
335

```

```

336     MPI_Finalize();
337     return 0;
338 }

```

## Appendix - Poisson solver

The parallel Poisson solver used the following code:

```

1  /*
2  * MPI_Poisson.c
3  * 2D Poisson equation solver (parallel version)
4  */
5
6  #include <stdio.h>
7  #include <stdlib.h>
8  #include <math.h>
9  #include <time.h>
10 #include <mpi.h>
11 #include <assert.h>
12
13 #define DEBUG 0
14
15 #define max(a,b) ((a)>(b)?a:b)
16
17 // defines for Exercises!
18
19 #define SOR 1
20 #define MONITOR_ERROR 1
21 #define MONITOR_ALLREDUCE 1
22 #define ALLREDUCE_COUNT 100
23
24 #define DEFINES_ON (SOR || MONITOR_ERROR || 0)
25 //defines end
26
27 enum
28 {
29     X_DIR, Y_DIR
30 };
31
32 // only needed for certain configs!
33 #ifdef SOR
34 double sor_omega = 1.9;
35 #endif
36 #ifdef MONITOR_ERROR
37 double *errors=NULL;
38 #endif
39 #ifdef MONITOR_ALLREDUCE
40 double all_reduce_time = 0;
41 #endif
42
43 /* global variables */
44 int gridsize[2];
45 double precision_goal; /* precision_goal of solution */
46 int max_iter; /* maximum number of iterations allowed */
47 int P; //total number of processes
48 int P_grid[2]; // process grid dimensions
49 MPI_Comm grid_comm; //grid communicator
50 MPI_Status status;
51 double h;
52
53 /* process specific globals*/
54 int proc_rank;
55 double wtime;
56 int proc_coord[2]; // coords of current process in processgrid
57 int proc_top, proc_right, proc_bottom, proc_left; // ranks of neighboring procs
58 // step 7
59 int offset[2] = {0,0};
60 // step 8
61 MPI_Datatype border_type[2];
62
63 /* benchmark related variables */

```



```

65 clock_t ticks;      /* number of systemticks */
66 int timer_on = 0;   /* is timer running? */
67
68 /* local grid related variables */
69 double **phi;       /* grid */
70 int **source;       /* TRUE if subgrid element is a source */
71 int dim[2];         /* grid dimensions */
72
73 void Setup_Grid();
74 double Do_Step(int parity);
75 void Solve();
76 void Write_Grid();
77 void Clean_Up();
78 void Debug(char *mesg, int terminate);
79 void start_timer();
80 void resume_timer();
81 void stop_timer();
82 void print_timer();
83
84 void start_timer()
85 {
86     if (!timer_on){
87         MPI_Barrier(grid_comm);
88         ticks = clock();
89         wtime = MPI_Wtime();
90         timer_on = 1;
91     }
92 }
93
94 void resume_timer()
95 {
96     if (!timer_on){
97         ticks = clock() - ticks;
98         wtime = MPI_Wtime() - wtime;
99         timer_on = 1;
100     }
101 }
102
103 void stop_timer()
104 {
105     if (timer_on){
106         ticks = clock() - ticks;
107         wtime = MPI_Wtime() - wtime;
108         timer_on = 0;
109     }
110 }
111
112 void print_timer()
113 {
114     if (timer_on){
115         stop_timer();
116         printf("(%) Elapsed Wtime %14.6f s (%5.1f%% CPU)\n", proc_rank, wtime, 100.0 * ticks
117             * (1.0 / CLOCKS_PER_SEC) / wtime);
118         resume_timer();
119     }
120     else{
121         printf("(%) Elapsed Wtime %14.6f s (%5.1f%% CPU)\n", proc_rank, wtime, 100.0 * ticks
122             * (1.0 / CLOCKS_PER_SEC) / wtime);
123     }
124 }
125
126 void Debug(char *mesg, int terminate)
127 {
128     if (DEBUG || terminate){
129         printf("%s\n", mesg);
130     }
131     if (terminate){
132         exit(1);
133     }
134 }
135
136 void Setup_Proc_Grid(int argc, char **argv){
137     int wrap_around[2];

```

```

136     int reorder;
137
138     Debug("My_MPI_Init",0);
139
140     // num of processes
141     MPI_Comm_size(MPI_COMM_WORLD, &P);
142
143     //calculate the number of processes per column and per row for the grid
144     if(argc>2){
145         P_grid[X_DIR] = atoi(argv[1]);
146         P_grid[Y_DIR] = atoi(argv[2]);
147         if(P_grid[X_DIR] * P_grid[Y_DIR] != P){
148             Debug("ERROR Proces grid dimensions do not match with P ", 1);
149         }
150         if (argc>3)
151         {
152             // get sor from args
153             sor_omega = atof(argv[3]);
154             printf("Set sor_omega over argv to %.4f\n", sor_omega);
155         }
156     }
157     else{
158         Debug("ERROR Wrong parameter input",1);
159     }
160
161     // Create process topology (2D grid)
162     wrap_around[X_DIR] = 0;
163     wrap_around[Y_DIR] = 0;
164     reorder = 1; //reorder process ranks
165
166     // create grid_comm
167     int ret = MPI_Cart_create(MPI_COMM_WORLD, 2, P_grid, wrap_around, reorder, &grid_comm);
168     if (ret != MPI_SUCCESS){
169         Debug("ERROR: MPI_Cart_create failed",1);
170     }
171     //get new rank and cartesian coords of this proc
172     MPI_Comm_rank(grid_comm, &proc_rank);
173     MPI_Cart_coords(grid_comm, proc_rank, 2, proc_coord);
174     printf("(%) (x,y)=(%,%)\\n", proc_rank, proc_coord[X_DIR], proc_coord[Y_DIR]);
175     //calc neighbours
176     // MPI_Cart_shift(grid_comm, Y_DIR, 1, &proc_bottom, &proc_top);
177     MPI_Cart_shift(grid_comm, Y_DIR, 1, &proc_top, &proc_bottom);
178     MPI_Cart_shift(grid_comm, X_DIR, 1, &proc_left, &proc_right);
179     printf("(%) top %, right %, bottom %, left %\\n", proc_rank, proc_top,
180     proc_right, proc_bottom, proc_left);
181 }
182 void Setup_Grid()
183 {
184     int x, y, s;
185     double source_x, source_y, source_val;
186     FILE *f;
187
188     Debug("Setup_Subgrid", 0);
189
190     if(proc_rank == 0){
191         f = fopen("input.dat", "r");
192         if (f == NULL){
193             Debug("Error opening input.dat", 1);
194         }
195         fscanf(f, "nx: %i\\n", &gridsize[X_DIR]);
196         fscanf(f, "ny: %i\\n", &gridsize[Y_DIR]);
197         fscanf(f, "precision goal: %lf\\n", &precision_goal);
198         fscanf(f, "max iterations: %i\\n", &max_iter);
199     }
200     MPI_Bcast(&gridsize, 2, MPI_INT, 0, grid_comm);
201     MPI_Bcast(&precision_goal, 1, MPI_DOUBLE, 0, grid_comm);
202     MPI_Bcast(&max_iter, 1, MPI_INT, 0, grid_comm);
203     h = 1;
204     /* Calculate dimensions of local subgrid */ //! We do that later now!
205     // dim[X_DIR] = gridsize[X_DIR] + 2;
206     // dim[Y_DIR] = gridsize[Y_DIR] + 2;
207

```

```

208  /// Step 7
209  int upper_offset[2] = {0,0};
210  // Calculate top left corner coordinates of local grid
211  offset[X_DIR] = gridsize[X_DIR] * proc_coord[X_DIR] / P_grid[X_DIR];
212  offset[Y_DIR] = gridsize[Y_DIR] * proc_coord[Y_DIR] / P_grid[Y_DIR];
213  upper_offset[X_DIR] = gridsize[X_DIR] * (proc_coord[X_DIR] + 1) / P_grid[X_DIR];
214  upper_offset[Y_DIR] = gridsize[Y_DIR] * (proc_coord[Y_DIR] + 1) / P_grid[Y_DIR];
215
216  // dimensions of local grid
217  dim[X_DIR] = upper_offset[X_DIR] - offset[X_DIR];
218  dim[Y_DIR] = upper_offset[Y_DIR] - offset[Y_DIR];
219  // Add space for rows/columns of neighboring grid
220  dim[X_DIR] += 2;
221  dim[Y_DIR] += 2;
222  /// Step 7 end
223
224  /* allocate memory */
225  if ((phi = malloc(dim[X_DIR] * sizeof(*phi))) == NULL){
226      Debug("Setup_Subgrid : malloc(phi) failed", 1);
227  }
228  if ((source = malloc(dim[X_DIR] * sizeof(*source))) == NULL){
229      Debug("Setup_Subgrid : malloc(source) failed", 1);
230  }
231  if ((phi[0] = malloc(dim[Y_DIR] * dim[X_DIR] * sizeof(**phi))) == NULL){
232      Debug("Setup_Subgrid : malloc(*phi) failed", 1);
233  }
234  if ((source[0] = malloc(dim[Y_DIR] * dim[X_DIR] * sizeof(**source))) == NULL){
235      Debug("Setup_Subgrid : malloc(*source) failed", 1);
236  }
237  for (x = 1; x < dim[X_DIR]; x++)
238  {
239      phi[x] = phi[0] + x * dim[Y_DIR];
240      source[x] = source[0] + x * dim[Y_DIR];
241  }
242
243  /* set all values to '0' */
244  for (x = 0; x < dim[X_DIR]; x++){
245      for (y = 0; y < dim[Y_DIR]; y++)
246      {
247          phi[x][y] = 0.0;
248          source[x][y] = 0;
249      }
250  }
251  /* put sources in field */
252  do{
253      if (proc_rank==0)
254      {
255          s = fscanf(f, "source: %lf %lf %lf\n", &source_x, &source_y, &source_val);
256      }
257      MPI_Bcast(&s, 1, MPI_INT, 0, grid_comm);
258      if (s==3){
259          MPI_Bcast(&source_x, 1, MPI_DOUBLE, 0, grid_comm);
260          MPI_Bcast(&source_y, 1, MPI_DOUBLE, 0, grid_comm);
261          MPI_Bcast(&source_val, 1, MPI_DOUBLE, 0, grid_comm);
262          x = source_x * gridsize[X_DIR];
263          y = source_y * gridsize[Y_DIR];
264          x = x + 1 - offset[X_DIR]; // Step 7 --> local grid transform
265          y = y + 1 - offset[Y_DIR]; // Step 7 --> local grid transform
266          if(x > 0 && x < dim[X_DIR] - 1 && y > 0 && y < dim[Y_DIR] - 1){ // check if in local
grid
267              phi[x][y] = source_val;
268              source[x][y] = 1;
269          }
270      }
271  }
272  while (s==3);
273
274  if(proc_rank==0){
275      fclose(f);
276  }
277 }
278
279 void Setup_MPI_Datatypes()

```

```

280 {
281     Debug("Setup_MPI_Datatypes",0);
282
283     // vertical data exchange (Y_Dir)
284     MPI_Type_vector(dim[X_DIR] - 2, 1, dim[Y_DIR], MPI_DOUBLE, &border_type[Y_DIR]);
285     // horizontal data exchange (X_Dir)
286     MPI_Type_vector(dim[Y_DIR] - 2, 1, 1, MPI_DOUBLE, &border_type[X_DIR]);
287
288     MPI_Type_commit(&border_type[Y_DIR]);
289     MPI_Type_commit(&border_type[X_DIR]);
290 }
291
292 void Exchange_Borders()
293 {
294     Debug("Exchange_Borders",0);
295     // top direction
296     MPI_Sendrecv(&phi[1][1], 1, border_type[Y_DIR], proc_top, 0, &phi[1][dim[Y_DIR] - 1], 1,
297     border_type[Y_DIR], proc_bottom, 0, grid_comm, &status);
298     // bottom direction
299     MPI_Sendrecv(&phi[1][dim[Y_DIR] - 2], 1, border_type[Y_DIR], proc_bottom, 0, &phi[1][0],
300     1, border_type[Y_DIR], proc_top, 0, grid_comm, &status);
301     // left direction
302     MPI_Sendrecv(&phi[1][1], 1, border_type[X_DIR], proc_left, 0, &phi[dim[X_DIR]-1][1], 1,
303     border_type[X_DIR], proc_right, 0, grid_comm, &status);
304     // right direction
305     MPI_Sendrecv(&phi[dim[X_DIR]-2][1], 1, border_type[X_DIR], proc_right, 0, &phi[0][1], 1,
306     border_type[X_DIR], proc_left, 0, grid_comm, &status);
307 }
308
309 double Do_Step(int parity)
310 {
311     int x, y;
312     double old_phi, c_ij;
313     double max_err = 0.0;
314
315     /* calculate interior of grid */
316     for (x = 1; x < dim[X_DIR] - 1; x++){
317         for (y = 1; y < dim[Y_DIR] - 1; y++){
318             if ((x + offset[X_DIR] + y + offset[Y_DIR]) % 2 == parity && source[x][y] != 1){
319                 old_phi = phi[x][y];
320                 #ifndef SOR
321                 phi[x][y] = (phi[x + 1][y] + phi[x - 1][y] + phi[x][y + 1] + phi[x][y - 1]) *
322                 0.25;
323                 #endif
324                 #ifdef SOR //! I'm not quite sure about the h and source parts here
325                 c_ij = (phi[x + 1][y] + phi[x - 1][y] + phi[x][y + 1] + phi[x][y - 1] + h*h*
326                 source[x][y]) * 0.25 - phi[x][y];
327                 phi[x][y] += sor_omega*c_ij;
328                 #endif
329                 if (max_err < fabs(old_phi - phi[x][y])){
330                     max_err = fabs(old_phi - phi[x][y]);
331                 }
332             }
333         }
334     }
335
336     return max_err;
337 }
338
339 void Solve()
340 {
341     int count = 0;
342     double delta;
343     double global_delta;
344     double delta1, delta2;
345
346     Debug("Solve", 0);
347
348     /* give global_delta a higher value then precision_goal */
349     global_delta = 2 * precision_goal;
350
351     while (global_delta > precision_goal && count < max_iter)
352     {

```

```

347     Debug("Do_Step 0", 0);
348     delta1 = Do_Step(0);
349     Exchange_Borders();
350     Debug("Do_Step 1", 0);
351     delta2 = Do_Step(1);
352     Exchange_Borders();
353     delta = max(delta1, delta2);
354     #ifdef MONITOR_ALLREDUCE
355     double time_ = MPI_Wtime();
356     #endif
357     #ifdef ALLREDUCE_COUNT
358     if(count % ALLREDUCE_COUNT == 0){
359         MPI_Allreduce(&delta, &global_delta, 1, MPI_DOUBLE, MPI_MAX, grid_comm);
360     }
361     #endif
362     #ifndef ALLREDUCE_COUNT
363     MPI_Allreduce(&delta, &global_delta, 1, MPI_DOUBLE, MPI_MAX, grid_comm);
364     #endif
365     #ifdef MONITOR_ALLREDUCE
366     all_reduce_time += MPI_Wtime() - time_;
367     #endif
368     #ifdef MONITOR_ERROR
369     if (proc_rank == 0)
370     {
371         errors[count] = global_delta;
372     }
373     #endif
374     count++;
375 }
376
377 printf("(%i) Number of iterations : %i\n", proc_rank, count);
378 #ifdef MONITOR_ALLREDUCE
379 printf("(%i) Allreduce time: %14.6f\n", proc_rank, all_reduce_time);
380 #endif
381 }
382
383 double* get_Global_Grid()
384 {
385     Debug("get_Global_Grid", 0);
386     ///!! DEBUG only
387     for (size_t i = 0; i < dim[X_DIR]; i++)
388     {
389         for (size_t j = 0; j < dim[Y_DIR]; j++)
390         {
391             phi[i][j] = proc_rank;
392         }
393     }
394
395     // only process 0 needs to store all data!
396     double* global_phi = NULL;
397     if (proc_rank == 0) {
398         global_phi = malloc(gridsize[X_DIR] * gridsz[e[Y_DIR] * sizeof(double));
399         if (global_phi == NULL) {
400             Debug("get_Global_Grid : malloc(global_phi) failed", 1);
401         }
402     }
403
404     // copy own part into buffer - flatten!
405     size_t buf_size = (dim[X_DIR] - 2) * (dim[Y_DIR] - 2) * sizeof(double);
406     double* local_phi = malloc(buf_size);
407     int idx = 0;
408     for (int x = 1; x < dim[X_DIR] - 1; x++) {
409         for (int y = 1; y < dim[Y_DIR] - 1; y++) {
410             local_phi[idx++] = phi[x][y];
411         }
412     }
413
414     printf("I'm proc %d and i have a buffer of size %zu\n", proc_rank, buf_size);
415
416     // only proc 0 needs sendcounts and displacements for the gather operation
417     int* sendcounts = NULL;
418     int* displs = NULL;

```

```

420     if (proc_rank == 0) {
421         sendcounts = malloc(P * sizeof(int));
422         displs = malloc(P * sizeof(int));
423
424         // size and offset of different subgrids
425         //! Note that this only works if every process has the same subgrid
426         if (gridsize[X_DIR] % P_grid[X_DIR] != 0 || gridsize[Y_DIR] % P_grid[Y_DIR] != 0)
427         {
428             Debug("!!!A grid dimension is not a multiple of the P_grid in this direction!", 1)
429         };
430     }
431
432     int subgrid_width = gridsize[X_DIR] / P_grid[X_DIR];
433     int subgrid_height = gridsize[Y_DIR] / P_grid[Y_DIR];
434     for (int px = 0; px < P_grid[X_DIR]; px++) {
435         for (int py = 0; py < P_grid[Y_DIR]; py++) {
436             int rank = px * P_grid[Y_DIR] + py;
437             sendcounts[rank] = subgrid_width * subgrid_height;
438             displs[rank] = (px * subgrid_width * gridsize[Y_DIR]) + (py * subgrid_height);
439         }
440     }
441     Debug("get_Global_Grid : MPI_Gatherv", 0);
442     //! TODO this Gatherv does something wrong - all local grids are alright!!!
443     MPI_Gatherv(local_phi, (dim[X_DIR] - 2) * (dim[Y_DIR] - 2), MPI_DOUBLE, global_phi,
444     sendcounts, displs, MPI_DOUBLE, 0, MPI_COMM_WORLD);
445
446     free(local_phi);
447     if (proc_rank == 0) {
448         free(sendcounts);
449         free(displs);
450     }
451     return global_phi;
452 }
453
454 void Write_Grid_global(){
455     int x, y;
456     FILE *f;
457     char filename[40]; //seems dangerous to use a static buffer but let's go with the steps
458     sprintf(filename, "output_MPI_global%i.dat", proc_rank);
459     if ((f = fopen(filename, "w")) == NULL){
460         Debug("Write_Grid : fopen failed", 1);
461     }
462
463     Debug("Write_Grid", 0);
464
465     for (x = 1; x < dim[X_DIR]-1; x++){
466         for (y = 1; y < dim[Y_DIR]-1; y++){
467             int x_glob = x + offset[X_DIR];
468             int y_glob = y + offset[Y_DIR];
469             fprintf(f, "%i %i %f\n", x_glob, y_glob, phi[x][y]);
470         }
471     }
472     fclose(f);
473 }
474
475 void Write_Grid()
476 {
477     double* global_phi = get_Global_Grid();
478     if(proc_rank != 0){
479         assert (global_phi == NULL);
480         return;
481     }
482     int x, y;
483     FILE *f;
484     char filename[40]; //seems dangerous to use a static buffer but let's go with the steps
485     sprintf(filename, "output_MPI%i.dat", proc_rank);
486     if ((f = fopen(filename, "w")) == NULL){
487         Debug("Write_Grid : fopen failed", 1);
488     }
489
490     Debug("Write_Grid", 0);

```

```

491     for (x = 0; x < gridsize[X_DIR]; x++){
492         for (y = 0; y < gridsize[Y_DIR]; y++){
493             fprintf(f, "%i %i %f\n", x+1, y+1, global_phi[x*gridsize[Y_DIR] + y]);
494         }
495     }
496     fclose(f);
497     free(global_phi);
498 }
499
500 void Clean_Up()
501 {
502     Debug("Clean_Up", 0);
503
504     free(phi[0]);
505     free(phi);
506     free(source[0]);
507     free(source);
508     #ifdef MONITOR_ERROR
509     free(errors);
510     #endif
511 }
512
513 void setup_error_monitor(){
514     if (proc_rank != 0)
515     {
516         return;
517     }
518
519     errors = malloc(sizeof(double)*max_iter);
520 }
521 void write_errors(){
522     if(proc_rank != 0){
523         return;
524     }
525     FILE *f;
526     char filename[40]; //seems dangerous to use a static buffer but let's go with the steps
527     sprintf(filename, "errors_MPI.dat");
528     if ((f = fopen(filename, "w")) == NULL){
529         Debug("Write_Errors : fopen failed", 1);
530     }
531
532     Debug("Write_Errors", 0);
533
534     for (size_t i = 0; i < max_iter; ++i)
535     {
536         fprintf(f, "%f\n", errors[i]);
537     }
538     fclose(f);
539 }
540 int main(int argc, char **argv)
541 {
542     MPI_Init(&argc, &argv);
543     Setup_Proc_Grid(argc,argv); // was earlier MPI_Comm_rank(MPI_COMM_WORLD, &proc_rank);
544     start_timer();
545
546     Setup_Grid();
547     Setup_MPI_Datatypes();
548
549     #ifdef SOR
550     if (proc_rank == 0)
551     {
552         printf("SOR using omega: %.5f\n", sor_omega);
553     }
554     #endif
555     #ifdef MONITOR_ERROR
556     setup_error_monitor();
557     #endif
558
559     Solve();
560     #ifdef MONITOR_ERROR
561     write_errors();
562     #endif
563     // Write_Grid();

```

```
564     Write_Grid_global();
565     print_timer();
566
567     Clean_Up();
568     MPI_Finalize();
569     return 0;
570 }
```